



IN3050/IN4050 - Introduction to Artificial Intelligence and Machine Learning Optimization and Search

Ali Ramezani-Kebrya



Search / loss landscape



Optimization

We need

- A numerical representation x for all possible solutions to the problem
- A function $f(x)$ that tells us how good solution x is
- A way of finding
 - $\max f(x)$ if bigger $f(x)$ is better (benefit)
 - $\min f(x)$ if smaller $f(x)$ is better (cost)

how well do we predict given ground-truth?

Knapsack problem

TSP


Optimisation and Search

- **Continuous Optimization** is the mathematical discipline which is concerned with finding the maxima and minima of functions, possibly subject to constraints.
- **Discrete Optimization** is the activity of looking thoroughly in order to find an item with specified properties among a collection of items.



Discrete optimization

- No intermediate sol
Sol space is discrete*
- **Chip design**
 - Routing tracks during chip layout design
 - **Timetabling**
 - E.g.: Find a course time table with the minimum number of clashes for registered students
 - **Travelling salesman problem**
 - Optimization of travel routes similar logistics problems



Example: Travelling Salesman Problem (TSP)

- Given the coordinates of n cities, find the ***shortest closed tour*** which visits each ***once and only once*** (i.e. exactly once).
- Constraint :
 - all cities be visited, once and only once.



Some Optimization Methods

1. Exhaustive search
 2. Greedy search and hill climbing
 3. Simulated annealing
 4. Gradient descent/ascent
 - Not applicable for discrete optimization
-
- Handwritten annotations on the right side of the slide:
- A red curly brace groups the first three methods (1, 2, 3) under the word "Discrete".
 - A red curly brace groups the last method (4) under the word "Continuous".
 - A yellow arrow points upwards from the "Continuous" brace towards the word "Divide".
 - Below the "Divide" word, the text "to buckets" is written in yellow.


1. Exhaustive search (AKA brute-force search)

- Test all possible solutions, pick the best
- Guaranteed to find the optimal solution
- For TSP: Try every possible ordering of the cities.
Need to evaluate $N!$ different solutions
 - For 70 cities, $N! > 10^{100}$. That's more than the number of atoms in the universe.


Exhaustive search

Only works for simple discrete problems, but can be approximated in continuous problems

- Sample the space at regular intervals (grid search)
- Sample the space randomly N times



How can we be smarter than exhaustive search?




How can we be smarter than exhaustive search?

- Usually, search spaces have some local structure
- Similar solutions often have similar quality
- Making small changes to a solution, and measuring resulting quality, we can gradually move towards better solutions

2. Greedy search


- Only generates and evaluates a single solution
- Makes several locally optimal choices, hoping the result will be near a global optimum
- Details depend on the problem being solved



Video from poprythm, at Youtube:
<https://www.youtube.com/watch?v=SC5CX8drAtU>

Hill climbing

- Pick a solution as the current best (e.g. a random solution)
- Compare to neighbor solution(s) → How to find a neighbor?
 - If the neighbor is better, replace the current best
 - Repeat until we reach a certain number of evaluations




What is the disadvantage/limitation to hillclimbing?


stuck in local optima

Exploitation and Exploration

- Search methods should combine:
 - Trying completely new solutions (like in exhaustive search) => **Exploration**
 - Trying to improve the current best solution by local search => **Exploitation**



Example: n-armed bandit





0.8


1.2

0.5


1.6



Exhaustive search – pure exploration




Hill Climbing – pure exploitation



Global optimization

- Most of the time, we must expect the problem to have **many local optima**
- Ideally, we want to find the best local optimum: **the global optimum**
- The best strategy is often to **combine exploration and exploitation**





How can we combine exploration and exploitation?

Allow mis steps to explore
Simulated Annealing

Start from a number of random
initial solutions
Evolutionary Alg

Mixed solution



Local optima

Algorithms like greedy search, hill climbing and gradient ascent/descent can only find local optima:

- They will only move through a strictly improving chain of neighbors
- Once they find a solution with no better neighbors they stop

Going the wrong way

What if we modified the **hill climber** to sometimes choose worse solutions?

- Goal: avoid getting stuck in a local optimum
- Always keep the new solution if it is better
- However, if it is worse, we'd still want to keep it sometimes, i.e. with some probability

3. Annealing

A thermal process for obtaining low energy states of a solid in a heat bath:


- Increase the temperature of the heat bath to a the point at which the solid melts
- Decrease the temperature slowly
- If done slowly enough, the particles arrange themselves in the minimum energy state

Simulated annealing


- Set an initial temperature T
- Pick an initial solution
- Repeat:
 - Pick a solution neighboring the current solution
 - If the new one is better, keep it
 - Otherwise, keep the new one with probability p
 - p depends on the **difference in quality** and the **temperature**. high temp -> high p
(more randomness)
- Reduce T

high T : high exploration
low T : high exploitation

Exploration ↑



Simulated Annealing Illustrated



Continuous optimization

- **Mechanics**
 - Optimized design of mechanical shapes etc.
- **Economics**
 - Portfolio selection, pricing risk management etc.
- **Control engineering**
 - Process engineering, robotics etc.



4. Gradient ascent / descent

$$\underline{x} \in \mathbb{R}^n$$

In continuous optimization we may be able to calculate the gradient of $f(x)$:

$$f: \mathbb{R}^n \rightarrow \mathbb{R}$$

gradient : function value
change along all directions


$$\nabla f(\underline{x}) = \begin{bmatrix} \frac{\delta f(\underline{x})}{\delta x_0} \\ \frac{\delta f(\underline{x})}{\delta x_1} \\ \vdots \\ \frac{\delta f(\underline{x})}{\delta x_n} \end{bmatrix}$$

$$\begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix}$$

The gradient tells us in which direction $f(\underline{x})$ increases the most

4. Gradient ascent / descent

similar to
hill climbing
gradient
tells us
the direction



Gradient ascent / descent (subtract)

Starting from $x^{(0)}$, we can iteratively find higher $f(x^{(k+1)})$ by adding a value proportional to the gradient to $x^{(k)}$:

$$\text{GA: } x^{(k+1)} = x^{(k)} + \gamma \nabla f(x^{(k)})$$

$$\text{GD: } x^{(k+1)} \leftarrow x^{(k)} - \gamma \nabla f(x^{(k)})$$

\hookrightarrow step-size or learning rate

Gradient Descent: Algorithm

Start with a point (guess)


Repeat

Determine a descent direction

Choose a step

Update

Until stopping criterion is satisfied






FIGURE 9.3 *Left:* In an ideal world we would know how to go to the minimum directly. In practice, we don't, so we have to approximate it by something like *right:* moving in the direction of steepest descent at each stage.

Source: Marsland

Gradient Descent: Algorithm

γ : how fast we would move!

Start with a point (guess)


Repeat

Determine a descent direction

Choose a step (using gradient)

Update

Until stopping criterion is satisfied



Gradient Descent: Algorithm

Start with a point (guess)


Repeat

Determine a descent direction

Choose a step

Update

Until stopping criterion is satisfied



Gradient Descent: Algorithm

We adapt & tune γ

Start with a point (guess)


Repeat

Determine a descent direction

Choose a step

Update


Until stopping criterion is satisfied



Stop when “close”
from minimum




gradient = 0 & function is flat
to



“No Free Lunch” Theorem

- No search method is **best for all problems**
- Choose the method and **search operators** that suits your problem
- There are however some algorithms that aim to do well **across a range of problems**
 - Evolutionary algorithms are one example

Performance of methods on problems



Summary

- Two classes of problems in optimization:
 - Discrete and Continuous
- Optimization methods:
 - Exhaustive search,
 - Greedy search
 - hill climbing
 - simulated annealing
 - gradient descent
- Exploration vs exploitation