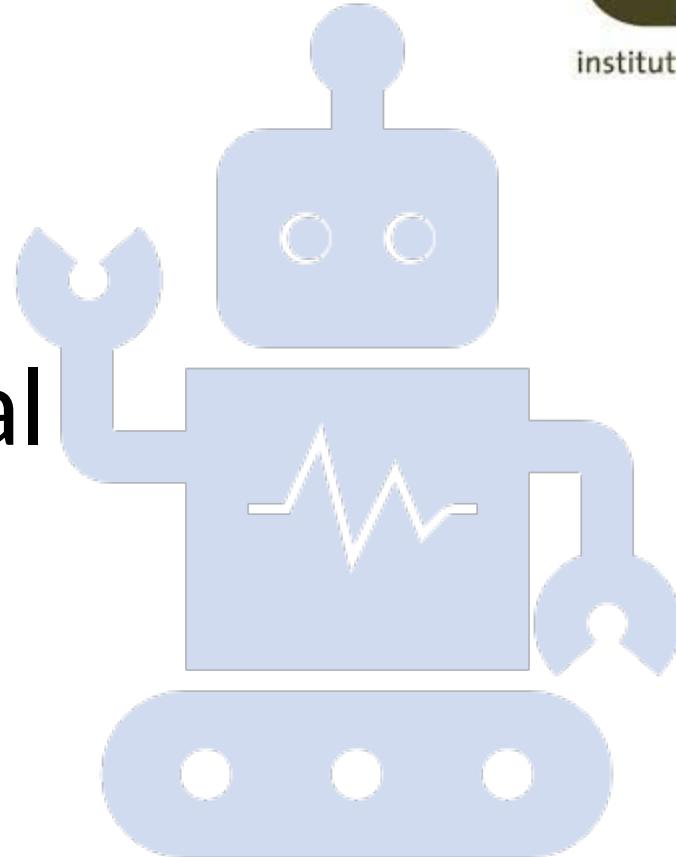


## IN3050/IN4050 - Introduction to Artificial Intelligence and Machine Learning

Lecture 3

Supervised learning 1

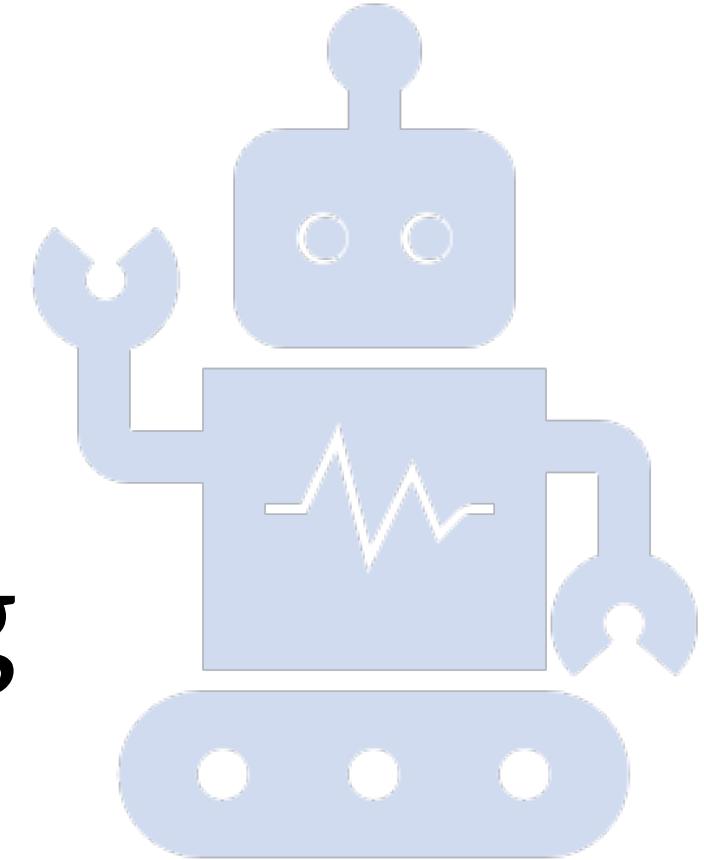
Ali Ramezani-Kebrya





# 5.1 Machine learning

IN3050/IN4050 Introduction to Artificial Intelligence  
and Machine Learning



Automated Learning of meaningful patterns in data (Understanding  
Machine Learning)

# Machine Learning

Converting experience in to expertise/knowledge

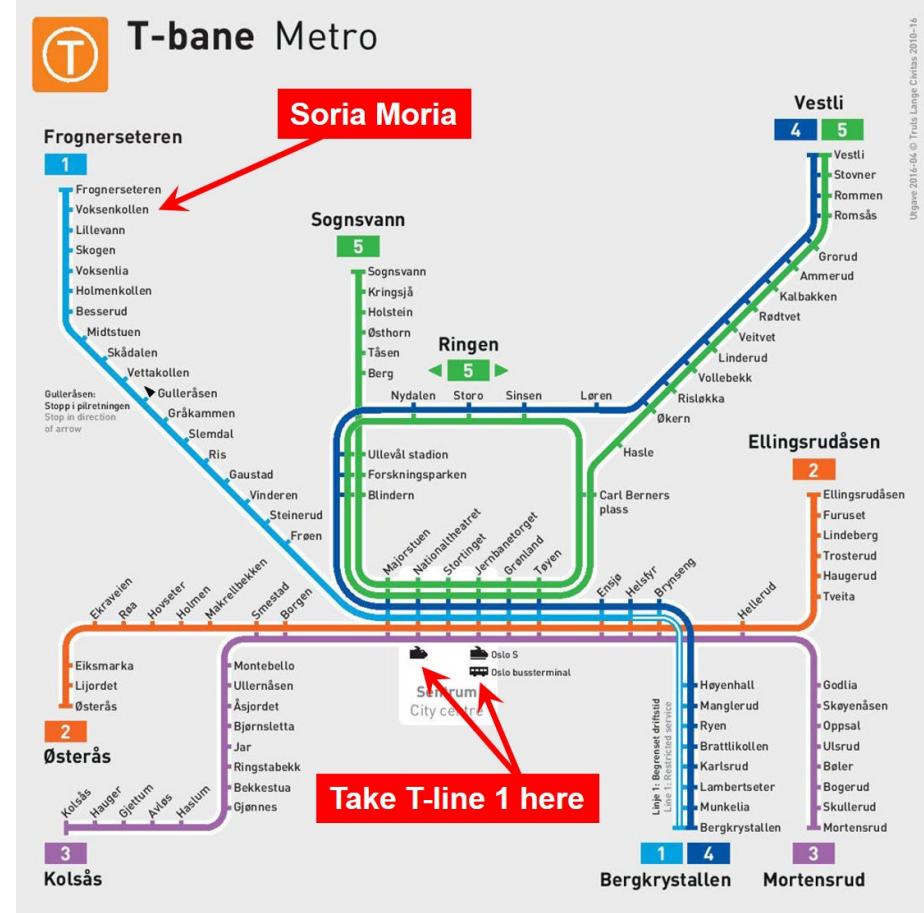
**Machine learning (ML)** is the study of computer algorithms that improve automatically through experience. (Wikipedia)

- (various more or less precise and more or less equivalent definitions)

# What does it mean to learn?

- If humans learn:
  - all the timetables for all the public transport in Oslo, would you say that they have learned?
  - all the entries in a bilingual dictionary?
- If a machine does the same, has it learned?

Generalization Matters  
(unseen data)



# Machine Learning

**Machine learning** (ML) is the study of computer algorithms that improve automatically through experience. (Wikipedia)

In particular :

- **Generalization:** Provide sensible outputs for inputs not encountered during training
- **extracting relevant information** from data and **applying it to analyze new data.**

# Formalizing the Learning Problem (Remember?)

- Learning = Improving with experience at some task
  - Improve over task  $T$
  - with respect to performance measure  $P$
  - based on experience  $E$

# Types of ML (Remember?)

## Supervised learning

- Learn from labeled data



- 



?

## Unsupervised learning

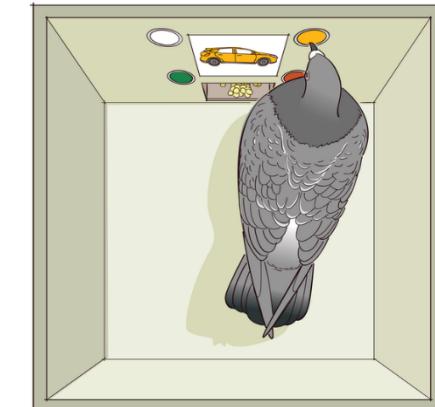
- No labeled data



- Task: identify similarities and categorize together

## Reinforcement learning

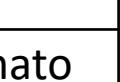
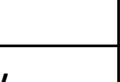
- Training with rewards (and punishments)

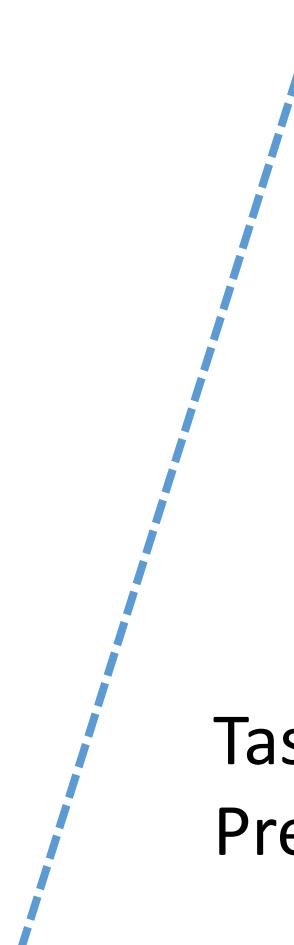


Source: Wikipedia

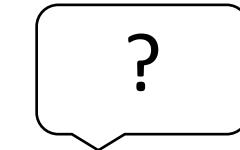
# Supervised learning

= Learning from exemplars

Training data							
							apple
							pear
							tomato
							cow
							dog
							horse
Each item is labelled							



Task:  
Predict the label on unseen items

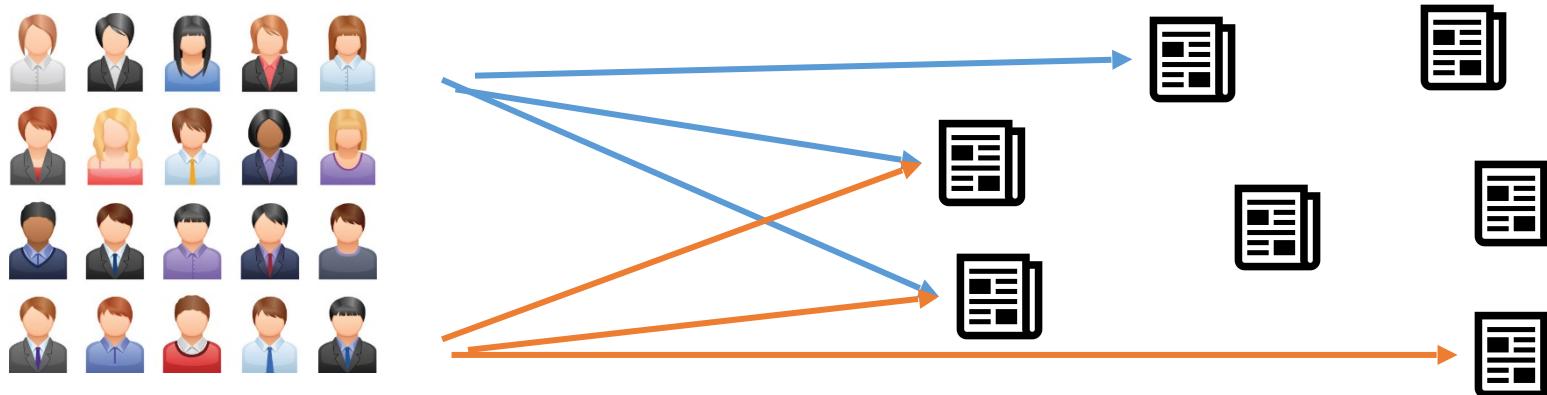


# Unsupervised learning

- *Can you sort the Lego bricks?*
  - (No instruction on how)
- You may choose sorting on
  - Color, or
  - Size, or
  - Shape, or
  - A combination
- I cannot know beforehand what you choose, but
- The result might be helpful

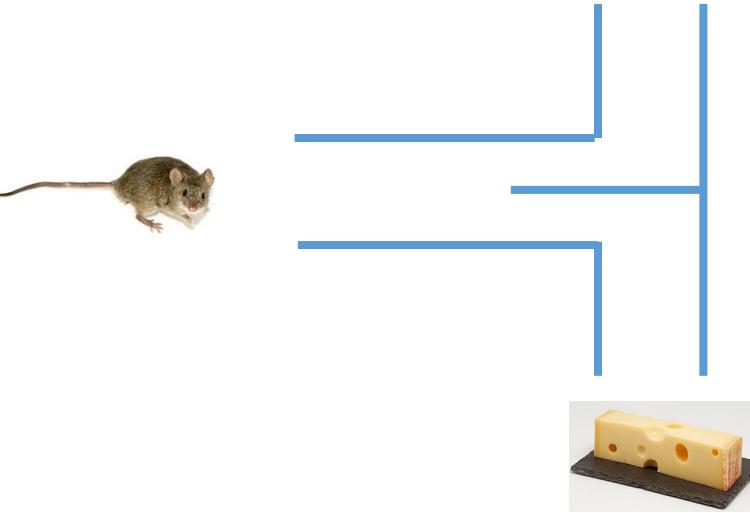
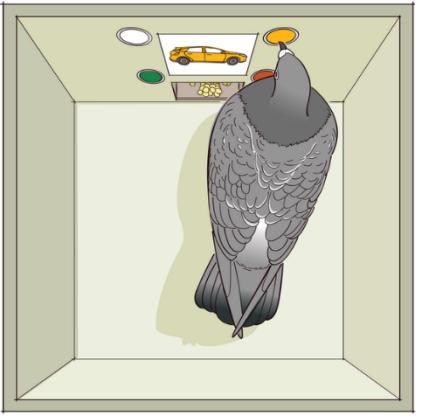


# Unsupervised learning, example 2



- Everybody (Facebook, Schibsted, ..) collects what you are reading
  - Readers who read the same stories are considered similar
  - Unsupervised learning can cluster readers based on what they read
  - Recommend new stories based on what others in your cluster have read
  - Convenient for you
  - Danger: Echo chambers
- population is divided into multiple cultures  
not reading others'/ interacting with others*

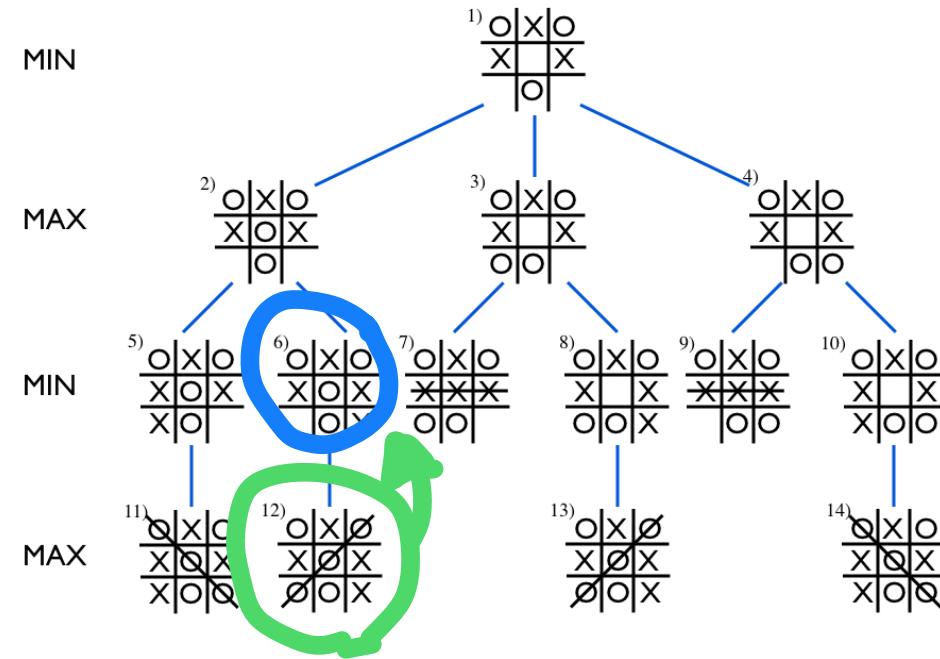
# Reinforcement learning (from Operant Conditioning in Psychology)



- The pigeon is given a reward if it pecks the right button when it sees the image.
- It learns this task
- An animal can learn the way around a maze if it gets a reward at the end and practices repeatedly
  - More complex mazes

# Game playing

- A game tree is like a maze
- By playing many times you may recognize first a winning position
- Next time you may recognize a position immediately before a winning position etc.



<https://materiaalit.github.io/intro-to-ai/part2/>

# AlphaZero

- It beats humans in Go and Chess
- Totally self-learned by playing against itself:
  - Reinforcement learning
  - Neural nets to generalize over game states
- Human plays, e.g., Carlsen, has learned new strategies from the program
  - “The *h* pawn”



# The next weeks: supervised learning

## Learning algorithms

1. Nearest neighbors
2. Perceptron
3. Linear regression
4. Simple neural networks
5. Multi-layer neural networks
  - Backpropagation
  - and more

## Aspects of the algorithms

- What are the underlying ideas?
- How are they implemented?
  - (1-5)
- What are the use cases?
- How can we apply them practically?

# Next weeks: The machine learning process

## Steps in the process

1. Data Collection and Preparation
2. Feature Selection and extraction
3. Algorithm Choice
4. Model Selection
5. Parameter Selection
6. Training
7. Evaluation
- Back and forth. 2-6

## What we will do

- Equally important as the algorithms.
- We will do this in parallel to the algorithms
- Starting with the basics and then introduce refinements

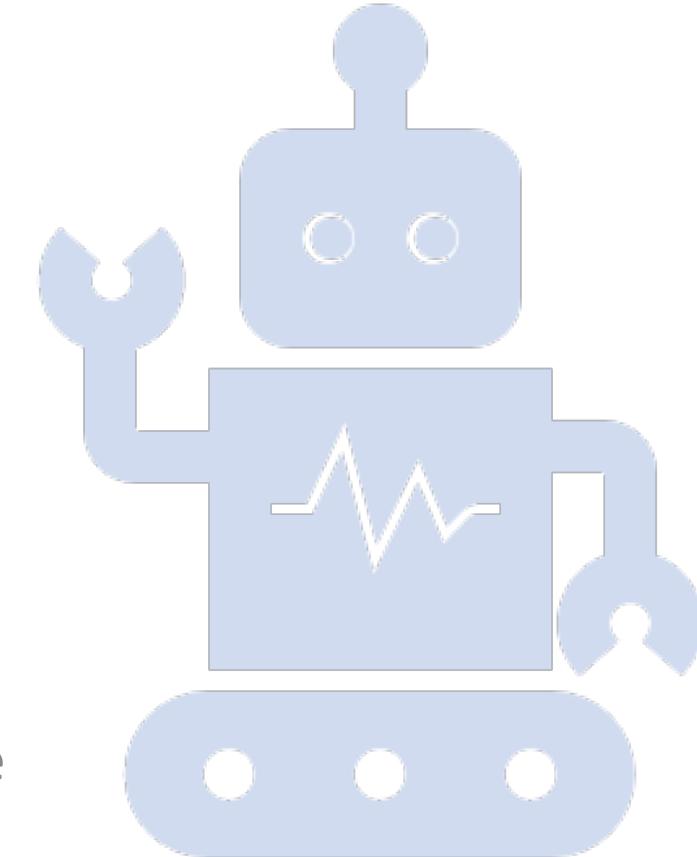
# Most and foremost supervised learning

- Techniques from supervised learning basis for the other two, e.g.,
  - Classification → clustering (e.g.,  $k$ -means clustering)
  - Recognizing similarities
  - Neural networks
- Experiments and evaluation are similar
  - Features
  - Test and training set
  - Evaluation measures
- Supervised learning is “the biggest field”
  - Most algorithms



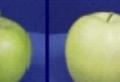
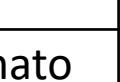
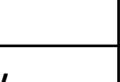
# 5.2 Classification and Features

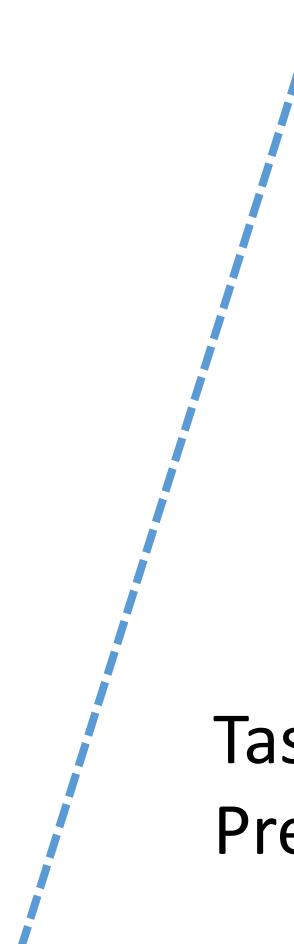
IN3050/IN4050 Introduction to Artificial Intelligence  
and Machine Learning



# Supervised learning

= Learning from exemplars

Training data							
							apple
							pear
							tomato
							cow
							dog
							horse
Each item is labeled							



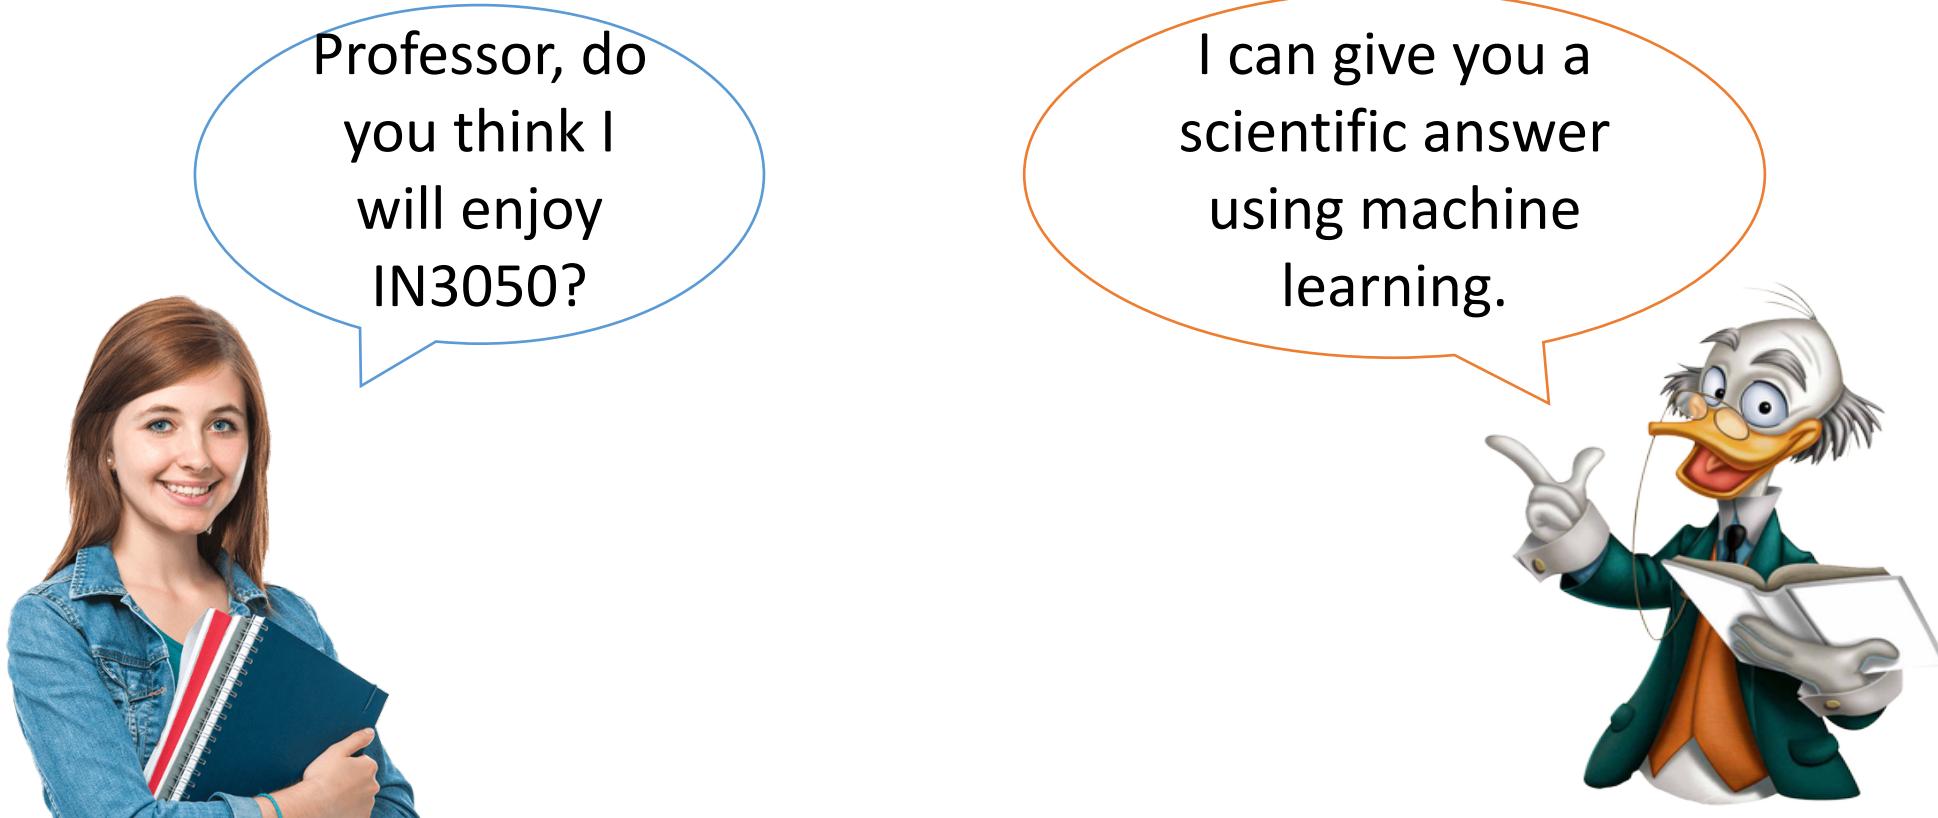
Task:  
Predict the label on unseen items



# What is a classifier?

- A **domain** of **objects/input features** we are to classify
- **Labels**: A finite set of labels
- **Classifier**: A mapping which maps each object to a unique label

# Example 1: (Decision Trees)



# Example 1: (Decision Trees)

## Domain



## Labels

- Will enjoy the course (yes/1)
- Will not enjoy the course (no/0)

- (Prospective IN3050/IN4050)  
Students

# Example 2: Spam mail

Domain



- E-mails

Labels

- Spam: yes
- Spam: no

# Example 3: The MNIST data set

Domain	Labels
<p>• Hand-written digits</p> <p>0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 5 5 5 5 5 5 5 5 5 5 5 5 5 5 5 6 6 6 6 6 6 6 6 6 6 6 6 6 6 6 7 7 7 7 7 7 7 7 7 7 7 7 7 7 7 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 9 9 9 9 9 9 9 9 9 9 9 9 9 9 9</p>	<p>0 1 2 3 4 5 6 7 8 9</p>

- Task
- To each hand-written picture of a digit, predict the correct digit
- There are 10 different classes

[https://en.wikipedia.org/wiki/MNIST\\_database](https://en.wikipedia.org/wiki/MNIST_database)

# Example 4: Language model

*The cat is on the ...*

- Task: Predict the next word!
- Labels:
  - A vocabulary of words:
    - *aardvark, ..., mat, ..., roof, ...*
    - E.g., 100,000 words/labels
- Domain:
  - Finite sequences of words



- Properties:
  - 1 billion training instances
  - Supervised learning
  - But you do not have to hand-label the training data.

large training data

# Classes

- **Binary classification:**
  - Two classes
- **Multiclass classification:**
  - Three or more classes
- Observe:
  - Some learning algorithms are by nature binary (e.g., [the perceptron](#)) and have to be adapted to multiclass classification
  - Other algorithms naturally admit any number of classes (e.g., [decision tree](#))

# Features

To be able to classify objects, we have to make a representation of them:

1. Decide on a set of **attributes/features** we can observe
2. Decide on the set of **possible values for each attribute**
3. Extract the **values of the attributes** for each object



# Example 1: (Decision Trees)

## Domain



- (Prospective IN3050/IN4050) Students

## Features

A questionnaire:

- Do you like mathematics?
  - Yes
  - no
- Do you have programming experience?
  - None
  - some (1 or 2 courses)
  - good (= 3 or more courses)
- Have you taken advanced machine learning courses?
  - Yes
  - no

# Results of the 2020 survey: a data set

Cand no	Enjoy maths	Programming	Adv. ML	Enjoy
1	Y	Good	N	Y
2	Y	Some	N	Y
3	N	Good	Y	N
4	N	None	N	N
5	N	Good	N	Y
6	N	Good	Y	Y
....				

# Transforming the classification task

- The task of predicting from a student



yes/no

- Is transformed to the task of predicting from some features

(Maths: Yes, Programming: Good, Adv.ML: No)



yes/no

# Example 2: email spam

	spam	chars	lines breaks	'dollar' occurs. numbers	'winner' occurs?	format	number
1	no	21,705	551	0	no	html	small
2	no	7,011	183	0	no	html	big
3	yes	631	28	0	no	text	none
4	no	2,454	61	0	no	text	small
5	no	41,623	1088	9	no	html	small
...							
50	no	15,829	242	0	no	html	small

From OpenIntro Statistics  
Creative Commons license

There are more variables  
(attributes) in the data set

- Data are typically represented in a **table**
- Each **column** one attribute
- Each **row** an observation (n-tuple, vector)
- (cf. Data base)

# Transforming the classification task

- The task of predicting from an e-mail



yes/no

- is transformed to the task of predicting from some features

(Chars: 21,705, Lines: 551, 'dollar': 0,  
'winner': no, format: html, number: small)



yes/no

# The larger picture

- This is how data sets are presented in texts on statistics or machine learning.
- But in real life, you want to apply ML to new tasks, then there is a lot of work before you have a data set like that:
  1. Data Collection and Preparation
  2. Feature Selection and extraction
- And for supervised learning, in particular
  3. Label the data, e.g., whether an x-ray shows cancer

# In this course

- We will mainly follow the trend and use pre-made data sets
- Concentrate on the more algorithmic and experimental parts of ML
- Techniques for, e.g., preprocessing and feature selection is to a certain degree domain specific. Hence this is left to e.g., courses in
  - Natural Language Processing
  - Image Processing
  - Bio Informatics
  - Robotics

# Example 2: email spam

	spam	chars	lines breaks	'dollar' occurs. numbers	'winner' occurs?	format	number
1	no	21,705	551	0	no	html	small
2	no	7,011	183	0	no	html	big
3	yes	631	28	0	no	text	none
4	no	2,454	61	0	no	text	small
5	no	41,623	1088	9	no	html	small
...							
50	no	15,829	242	0	no	html	small

50 observations, rows

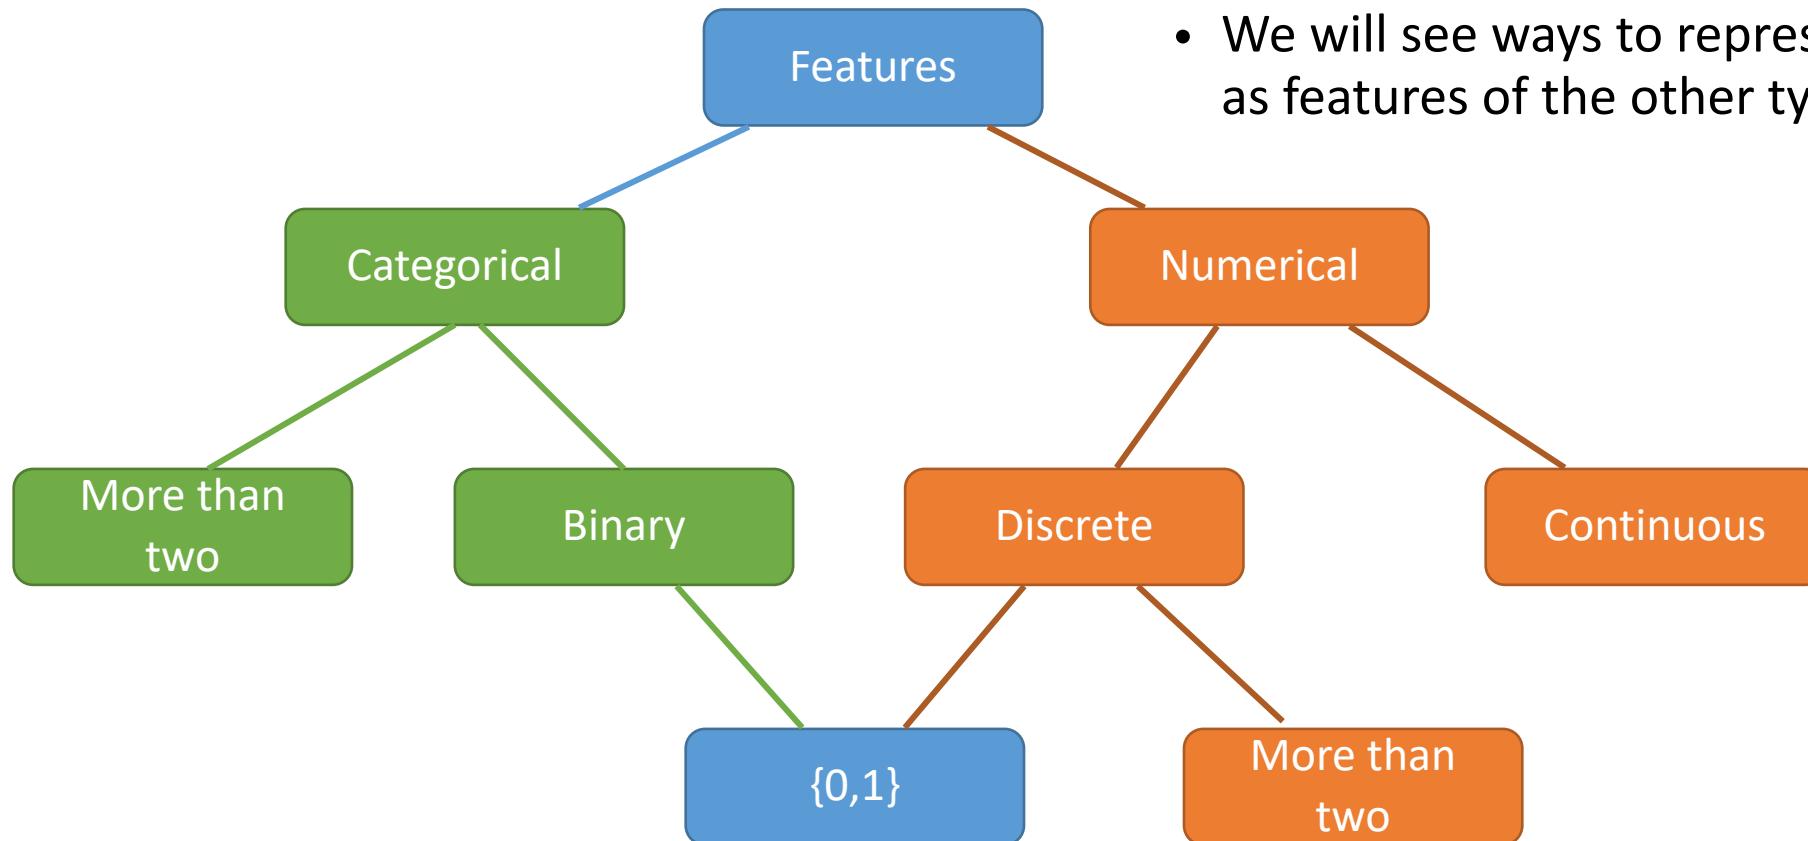
7 variables, columns

4 categorical variables

3 numerical variables

# Feature types

- Some algorithms can only take numerical features
  - $k$ NN, perceptron,
- Some algorithms take categorical features
  - Decision trees
- Some software can only take numerical features
  - Scikit learn
- We will see ways to represent a feature of one type as features of the other type



# From data set to features

- In some cases, the attributes of a data set can be used directly as features in an ML-algorithm.
- In other cases, it is necessary with some further preprocessing, e.g.
  - Transforming categorical data to numerical data
  - Scaling the data
  - Combining features
  - “Feature engineering”
  - Fixing the positions of the features
- The result is a fixed shape numerical vector

$$\mathbf{x} = (x_1, x_2, \dots, x_n)$$

# Transforming the classification task

- The task of predicting from an e-mail



yes/no

- is transformed to the task of predicting from some features

(Chars: 21,705, Lines: 551, 'dollar': 0,

'winner': no, format: html, number: small)



yes/no

- is transformed to the task of predicting from a numerical vector to a number

$x = (x_1, x_2, \dots, x_n)$



$y \in \{0, 1\}$

# Some words on notation

- The input as a numerical feature vector  $\mathbf{x} = (x_1, x_2, \dots, x_m)$ 
  - Bold face for vector
  - Sometimes start counting from 1, sometimes from 0
- Initially we consider classification as predicting a one-dimensional label:  $\mathbf{x} \rightarrow y$
- Marsland (ch. 2, p. 15) considers the output as a vector  $\mathbf{y} = (y_1, y_2, \dots, y_n)$
- We will use vector when the output has more than one dimension

# A word on notation

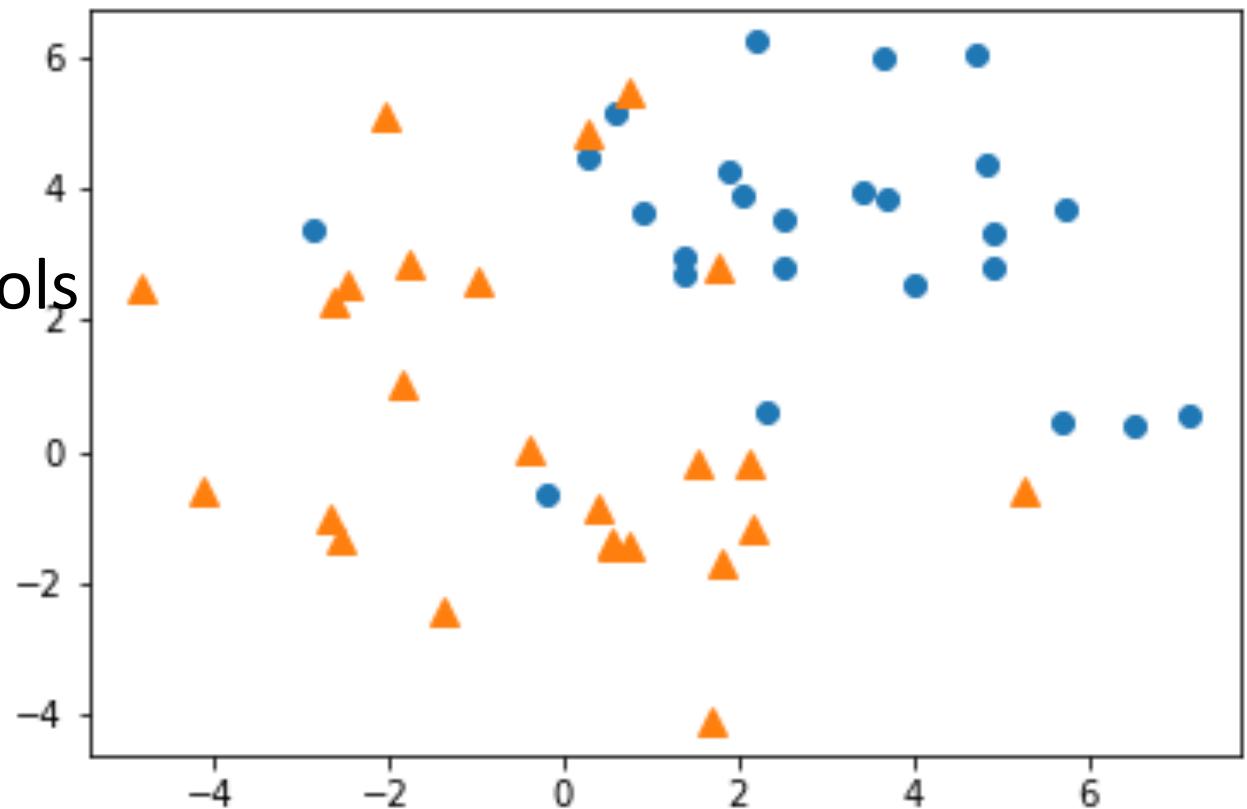
We will try to follow Marsland and use

- $\mathbf{x}_j = (x_{j,1}, x_{j,2}, \dots, x_{j,m})$  for input point number  $j$
- $t_j$  for a target value
- $y_j = f(\mathbf{x}_j)$  for a predicted value
- But beware that some use
  - $y_j$  for the target value and
  - $\hat{y}_j = f(\mathbf{x}_j)$  for the predicted value (or even  $\hat{y}_j = \hat{f}(\mathbf{x}_j)$ )

# ML algorithms and visualization

To illustrate ML-algorithms, it is easiest to use

- 2 numerical features
- show classes with colors, symbols (glyphs) or both



# Multi-label classification

- Classify an object with respect to several binary classes
- E.g., *Who is in the picture?*
- (Uncle Tom: 1, Aunt Mary: 1, Grandma: 0, Grandpa: 0, etc.)



# Supervised learning - two types

More on regression next week

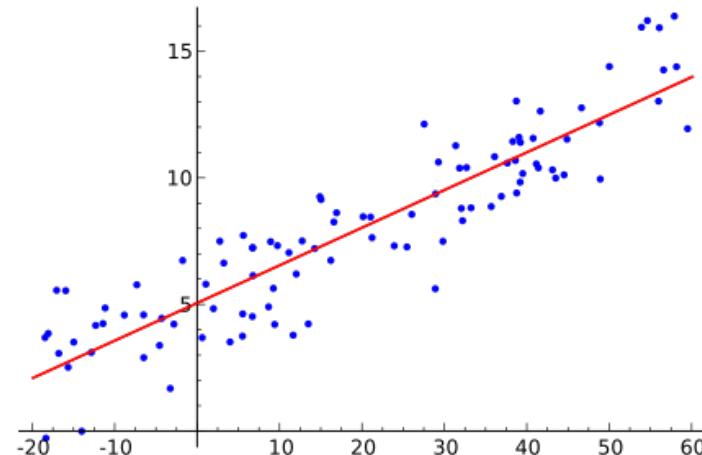
## Classification

- Assign a label (class) from a finite set of labels to an observation



## Regression

- Assign a numerical value to an observation
  - e.g., the temperature tomorrow



# Important concepts

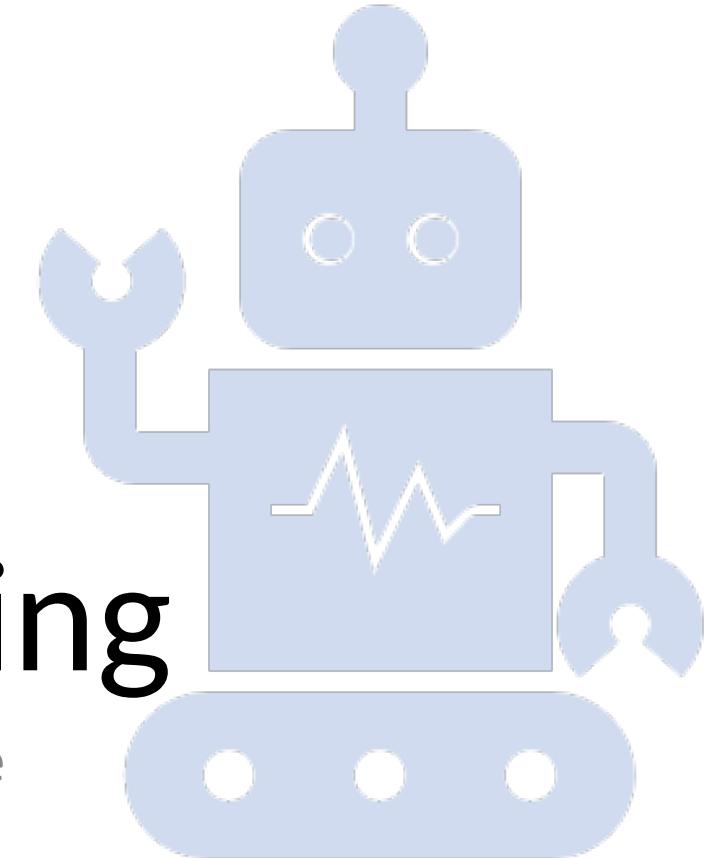
- Classifier, domain, label set
- Classifiers:
  - Binary
  - Multi-class classifier
- Multi-label classifiers
- Regression
- Features:
  - Categorical
  - Numerical





# 5.3 Supervised learning

IN3050/IN4050 Introduction to Artificial Intelligence  
and Machine Learning



# Supervised Learning

- Structure:
  - a well-defined set of observations,  $O$  (possible inputs)
  - a set of label values,  $L$  (possible output values)
- Goal: determine a mapping  $\gamma$ , from  $O$  to  $L$
- A training set of examples from  $O \times L$ ,  $\left\{ (o_1, t_1), (o_2, t_2), \dots, (o_n, t_n) \right\}$
- The **training** phase: Learn  $\gamma$  from the training set
  - Ideally,  $\gamma(o_i) = t_i$  on the training data. = the "supervision" of the learning
  - When  $\gamma$  is learned, it can be used to **predict** values for new items,  $\gamma(o)$ .

# Supervised Learning

1. **Feature extraction:** We first decide on the features, their possible values, extract them from the observations, and replace each observation with its features, e.g.,

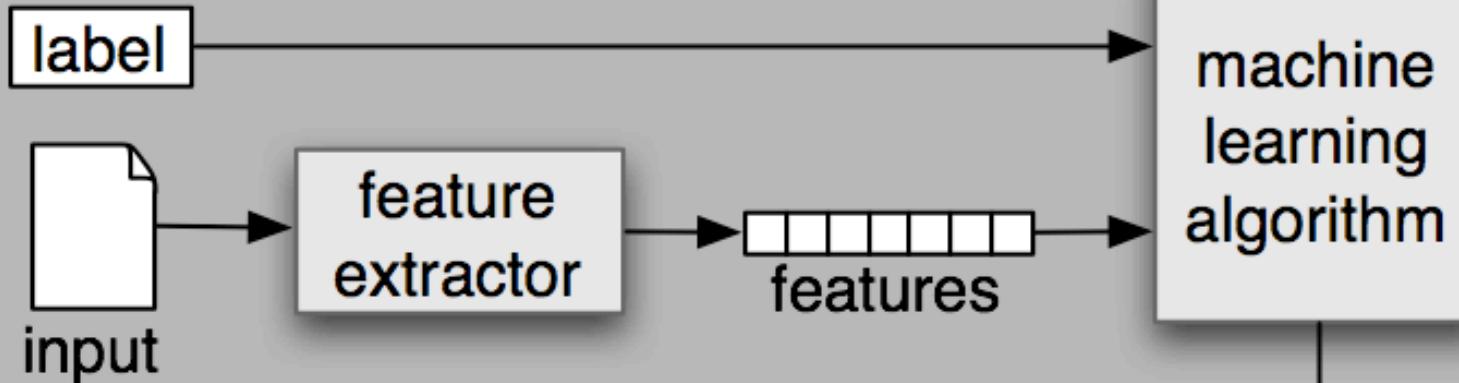
$o_j$  gets represented by  $\mathbf{x}_j = (x_{j,1}, x_{j,2}, \dots, x_{j,m})$

- Revised goal: determine a mapping  $f$ , from the set of features to L
- A training set of examples from  $\{(\mathbf{x}_1, t_1), (\mathbf{x}_2, t_2), \dots, (\mathbf{x}_n, t_n)\}$

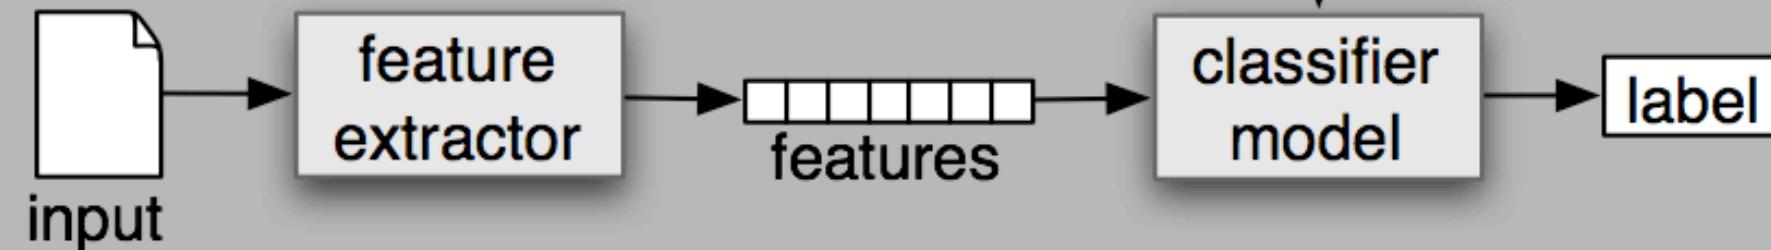
2. The **training** phase: Learn  $f$  from the training set
  - Ideally,  $f(\mathbf{x}_j) = t_j$  on the training data. = the "supervision" of the learning
3. When  $f$  is learned, it can be used to **predict** values for new items,  $f(\mathbf{x}')$

# Classification

## (a) Training

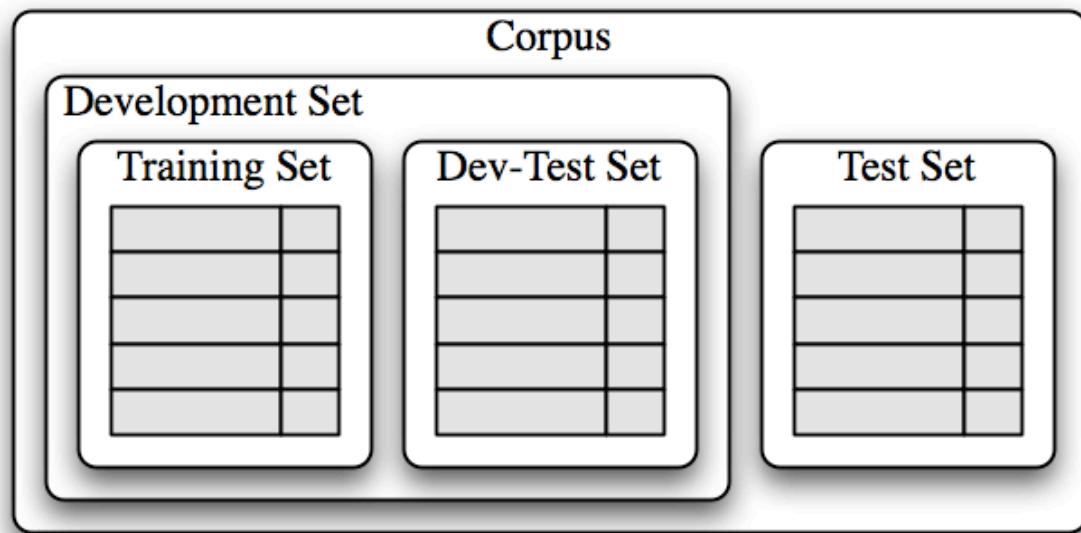


## (b) Prediction



# Training and test sets

- To measure improvement, we need (at least) two disjoint labeled sets:
  - Training set
  - Test set
- Train on the training set.
- Predict labels on the test set (after removing the labels)
- Compare the prediction to the given labels



<https://www.nltk.org/book/ch06.html>

- For repeated development we need (at least) two test sets.
  - One for repeated testing during development
  - One for final testing

# Confusion matrix and accuracy

Goal: Evaluate our spam classifier

- We run the classifier on the labeled test set (without the labels)
- Compare the predicted labels to the example labels and count
- We can present the numbers in a **confusion table**

		True label	
		Yes	NO
Predicted label	Yes	tp=150	fp=50
	No	fn=100	tn=200

- True positives, tp=150
- False positives, fp=50
- False negatives, fn=100
- True negatives, tn=200
- **Accuracy:**  
$$(tp+tn)/N = 350/500 = 0.7$$
  - (Marsland (2.2) p.23 is wrong!)

# More than two classes

		True label		
		spam	normal	urgent
Predicted label	spam	150	49	1
	normal	31	250	19
	urgent	19	31	50

## Accuracy:

- (sum of the diagonal)/N

$$\bullet = \frac{\#\{y_i | y_i = t_i\}}{\#\{y_i\}} = \frac{450}{600} = 0.75$$

## Observe

- There is no consensus regarding what should be the columns and what should be the rows
- (Marsland p.22 does it differently from p. 23)

# Important concepts

Steps:

- Feature extraction
- Training on development data
- Predicting on new data

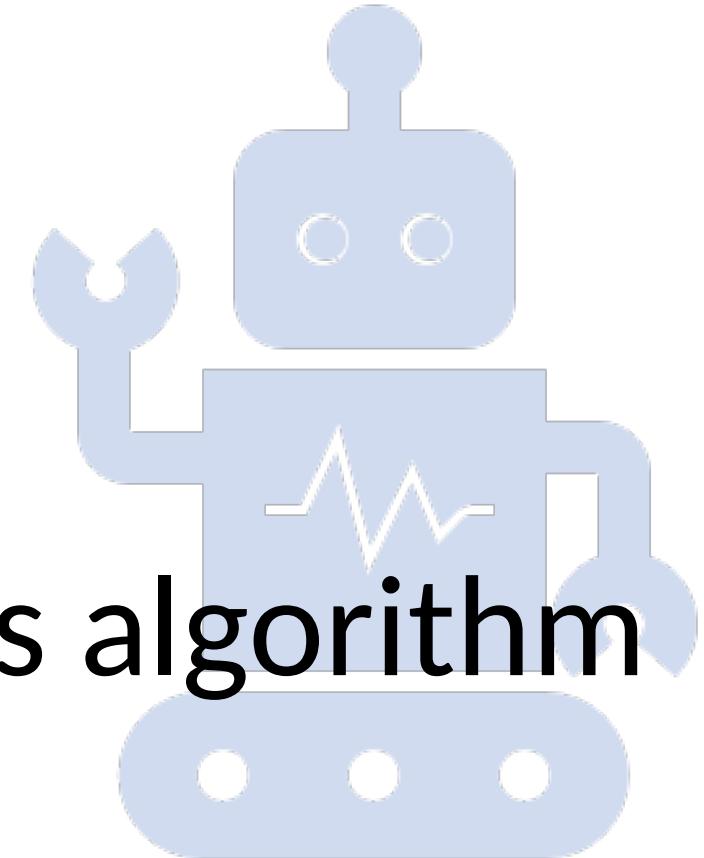
Split data:

- Development set:
  - Training set
  - Development test set
- Test set

Evaluation:

- Confusion matrix
- Accuracy





# 5.4 $k$ Nearest Neighbors algorithm

IN3050/IN4050 Introduction to Artificial Intelligence  
and Machine Learning

# Example:



Professor, do  
you think I  
will enjoy  
IN4050?

I can give you a  
scientific answer  
using machine  
learning.



# Another student



- Ask, say, 7 students who took the course whether they liked it or not.
- Some *yes*, some *no*-s
- I trust most the students who resemble me most, e.g.:
  - Taken the same courses as me
  - Like the same courses as I do
- I pick the 7 students who resemble me most, and ask them
- I trust the majority vote of this group

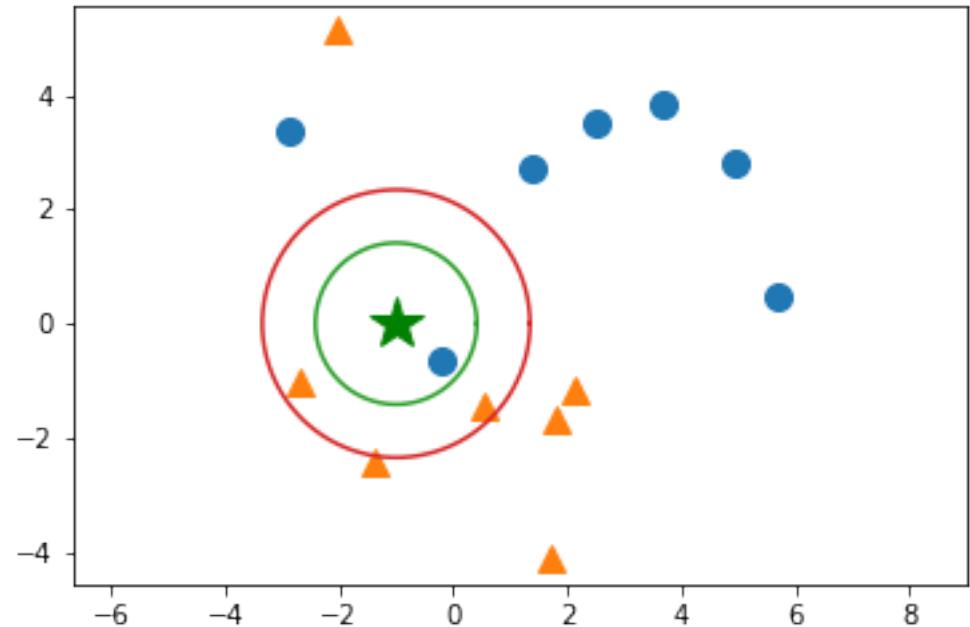
# $k$ Nearest Neighbors



- This is an example of  $k$  Nearest Neighbor algorithm, with  $k=7$
- Downside:
  - You have to consider every student who took the course last year to find the ones that resemble you most
  - Every student asking the question has to start the procedure from scratch

# kNN algorithm

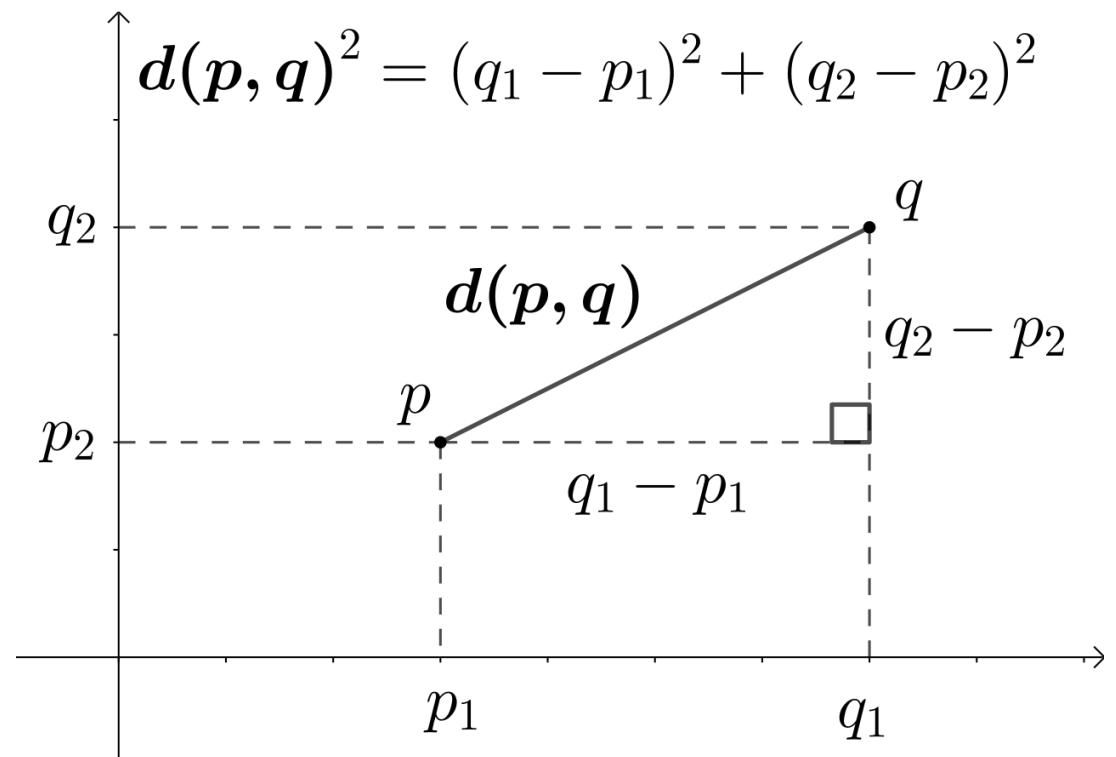
1. Calculate the distance to all the training instances
2. Pick the  $k$  nearest ones
3. Choose the majority class for this set



What is the result with  
 $k = 1?$   
 $k = 3?$

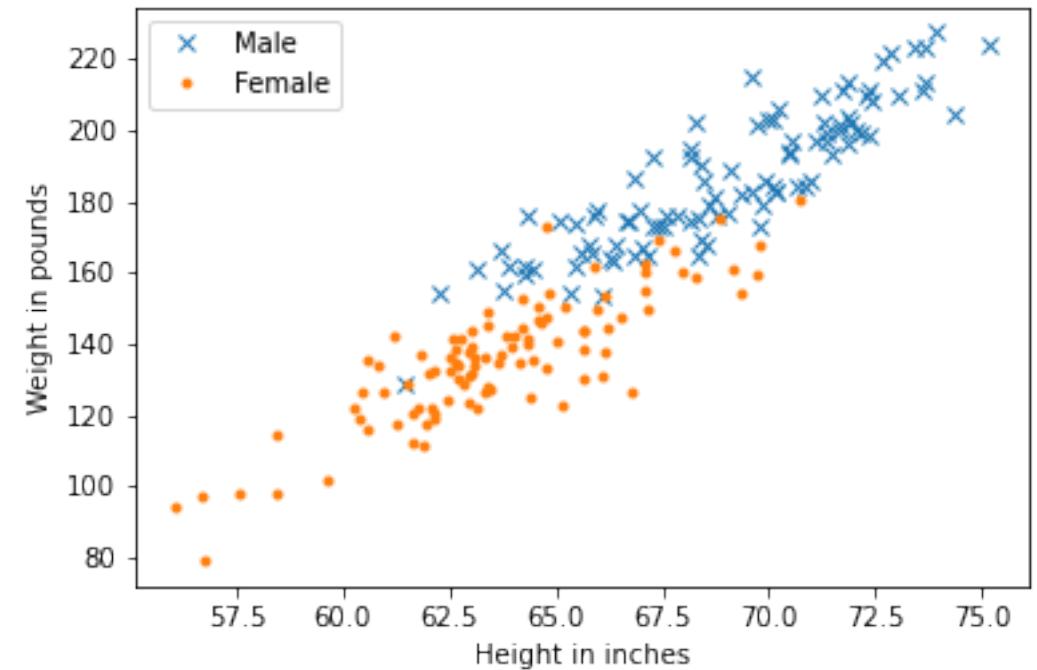
# Distance

- Assumption:
  - Points that are close together in feature space are similar
  - Some notion of distance in feature space, normally Euclidean distance (L2)



# Example: Height, weight, gender

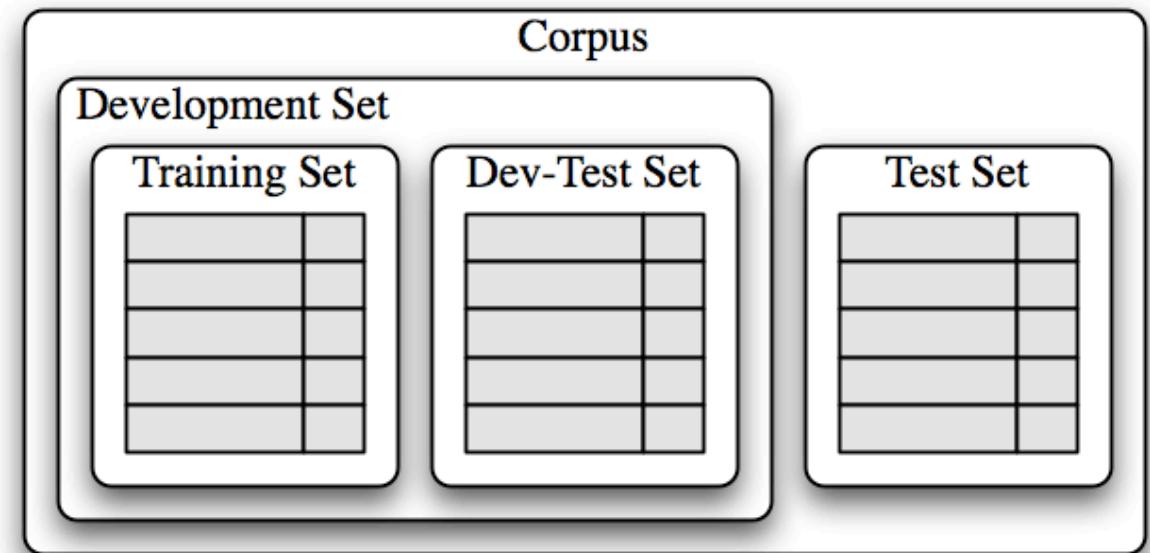
- Dataset:
  - 10,000 observations
  - 5,000 of each gender
  - Height in inches
  - Weight in pounds
- <https://www.kaggle.com/mustafaali96/weight-height>
- Processing:
  - Shuffled
  - Split:
    - 5000 for training
    - 2500 for development testing
    - 2500 for final testing



A random subset of 200 of the training data

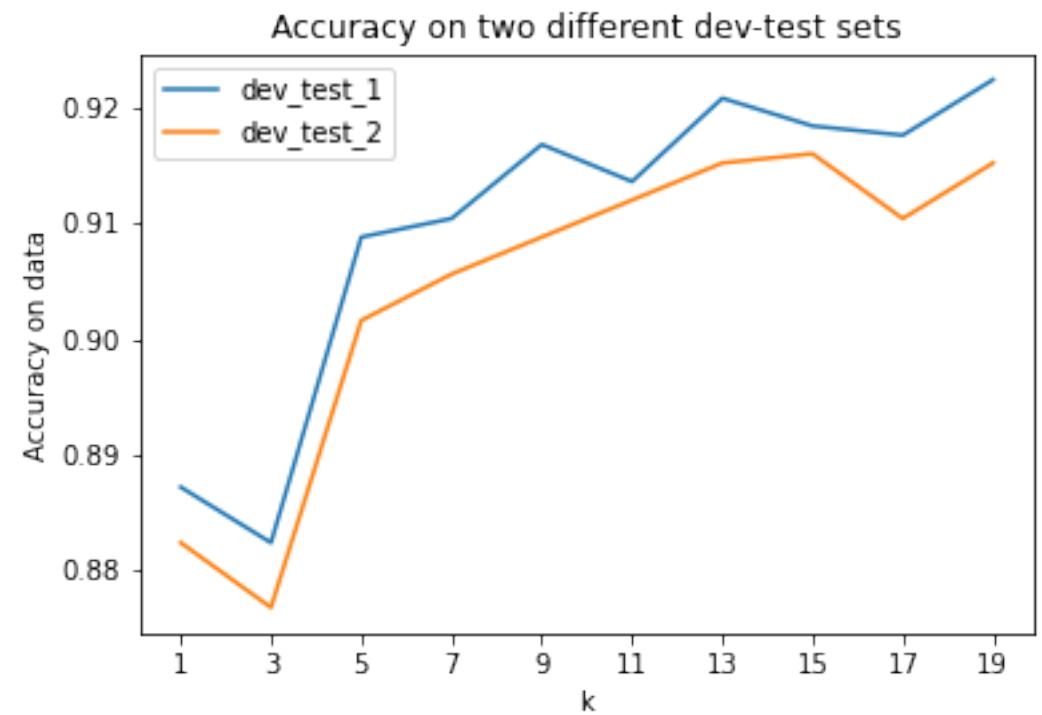
# Choosing $k$

- For alternative values of  $k$ :
  - Train on the training set
  - Evaluate on the dev-test set
- Choose the  $k$  which yields the best accuracy
- You may test with this  $k$  on the final test set.



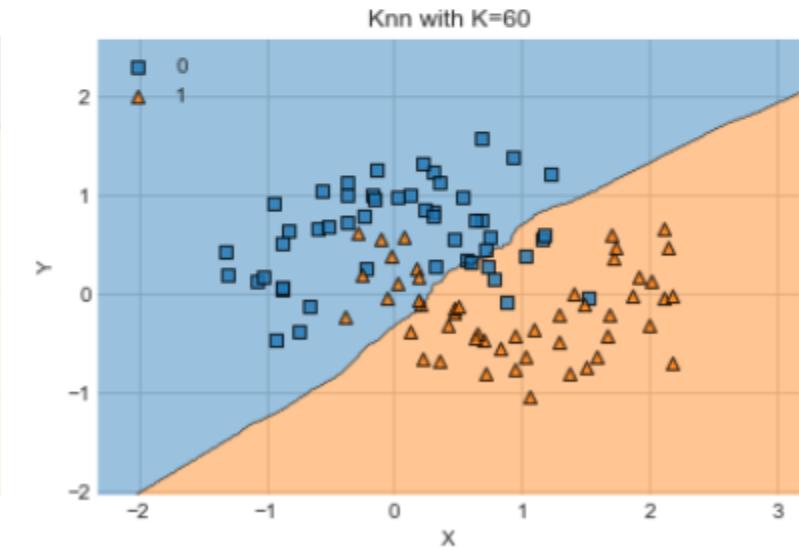
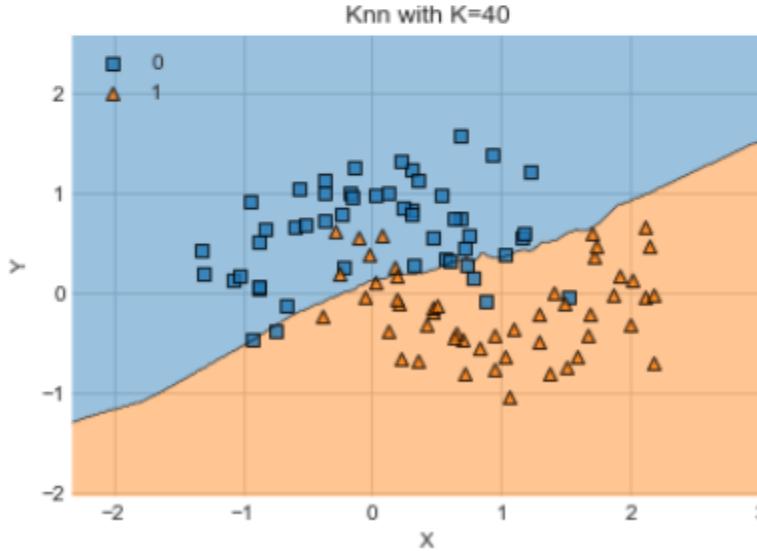
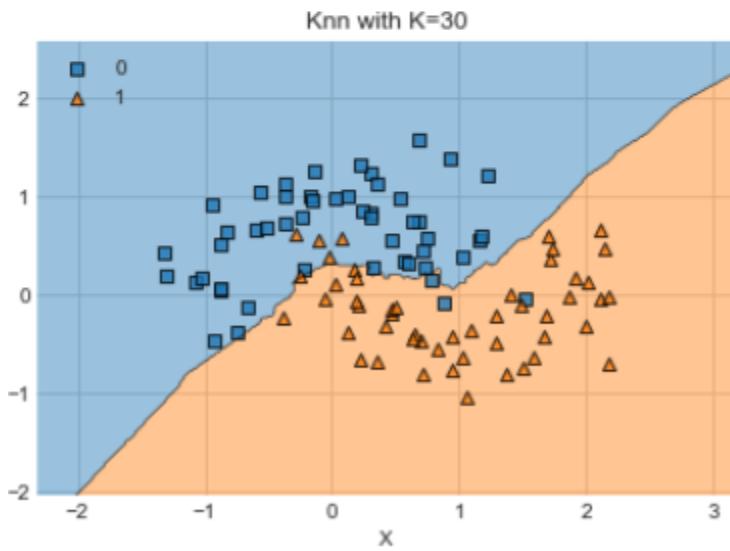
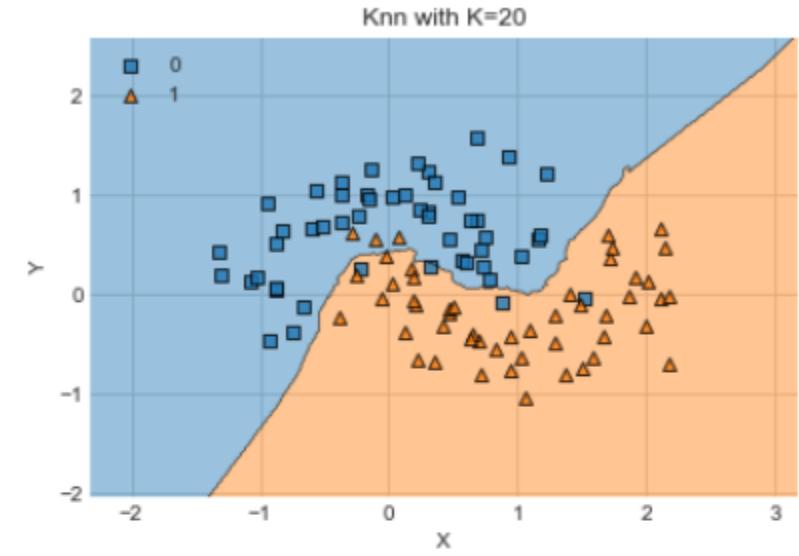
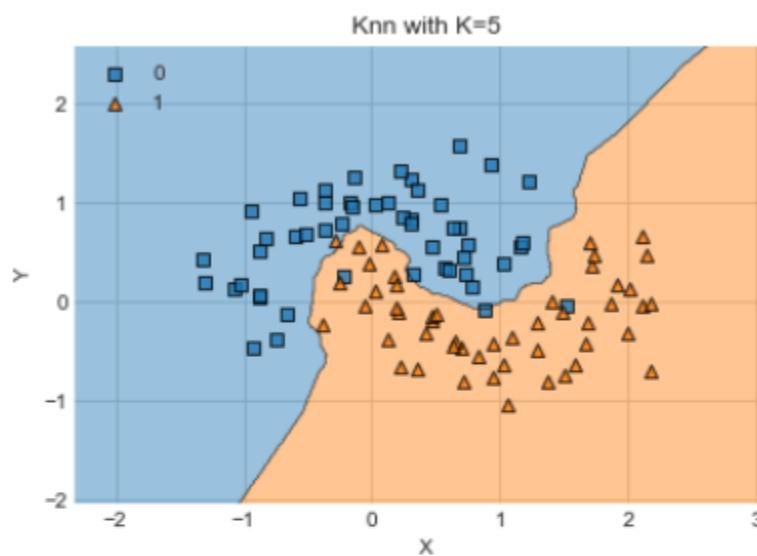
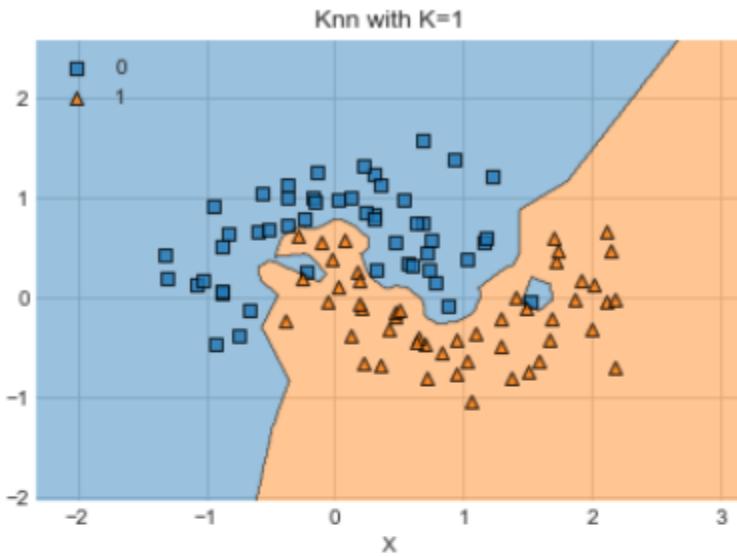
# $k$ NN with varying $k$

- Trained on the train set with various  $k$ 's
  - Accuracy on dev-test set
  - For this task: better with larger  $k$
- Split the dev-test randomly in two equal parts
  - Same tendency
  - Not the exact same numbers:
    - To be expected
  - Numbers like accuracy will vary with test set



# Is larger $k$ always better?

- Small  $k$ :
  - Good fit to training data
  - Danger of overfitting
- Larger  $k$ :
  - More general
- Next slide:
  - The squares and triangles are the true classes
  - The background colors show the decision boundary for the classifier with various  $k$ 's



KNN visualization for the U-shaped dataset

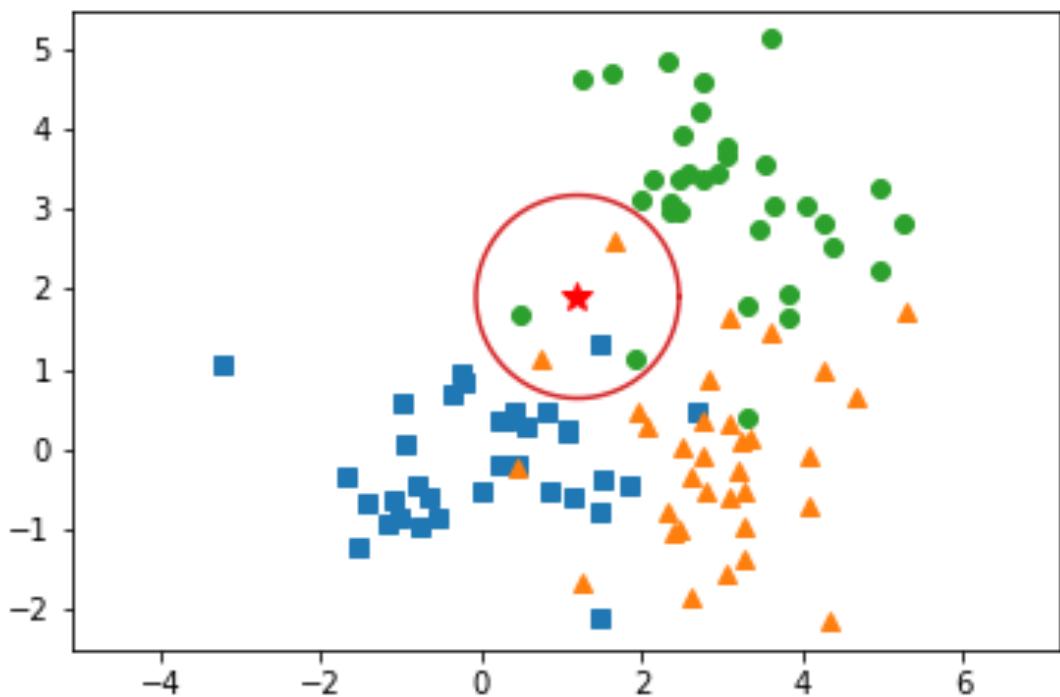
<https://towardsdatascience.com/knn-visualization-in-just-13-lines-of-code-32820d72c6b6>

# Properties of kNN

- Instance-based, no real training
  - (Fast to "train")
- Inefficient in predicting the label of new instances
  - Since it must consider all the training data
- One parameter:  $k$
- The distance measure may influence the result
- The scaling of the axes might influence the result:
  - Next!

# Footnote: More classes

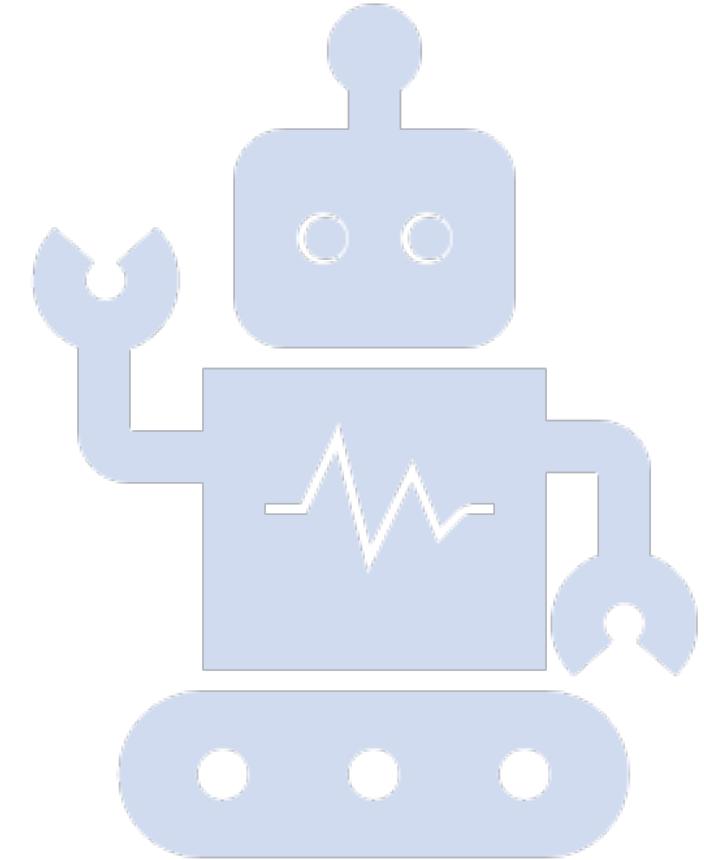
- A binary classifier with odd  $k$  always reaches a decision
- With more than 2 classes, there might be a draw
- Possible ways out
  - Alt. 1: Back-off to  $k-1$ , etc. until there is a majority class
  - Alt. 2: Weight points by inverse distance from target, take the weighted max.





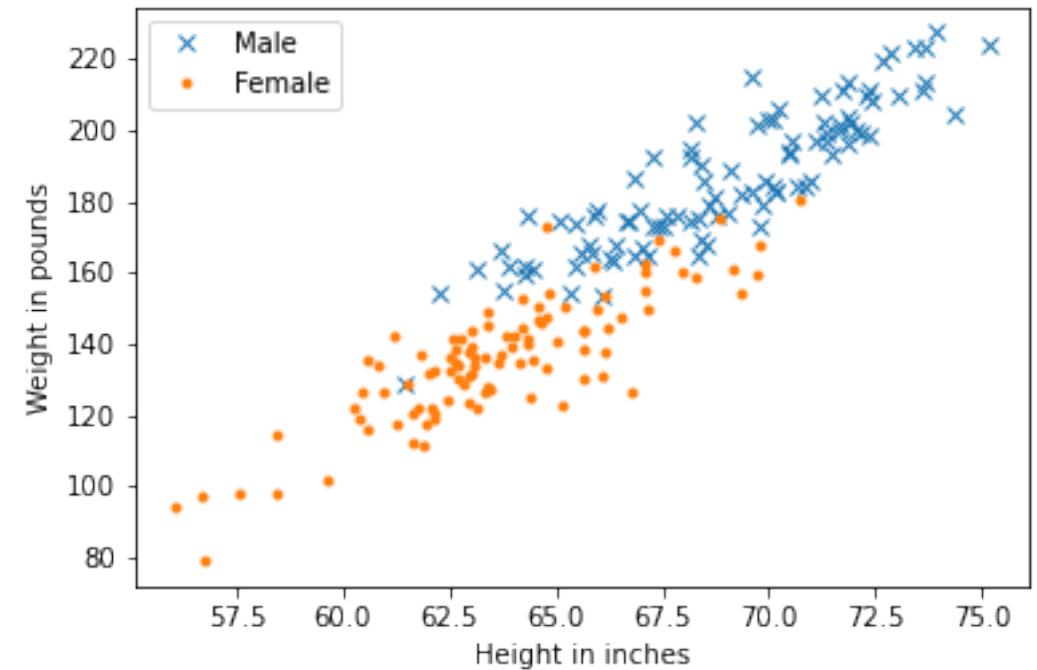
# 5.5 Feature scaling

IN3050/IN4050 Introduction to Artificial Intelligence  
and Machine Learning



# Example: Height, weight, gender

- Dataset:
  - 10,000 observations
  - 5,000 of each gender
  - Height in inches
  - Weight in pounds
- <https://www.kaggle.com/mustafaali96/weight-height>
- Processing:
  - Shuffled
  - Split:
    - 5000 for training
    - 2500 for development testing
    - 2500 for final testing



A random subset of 200 of the training data

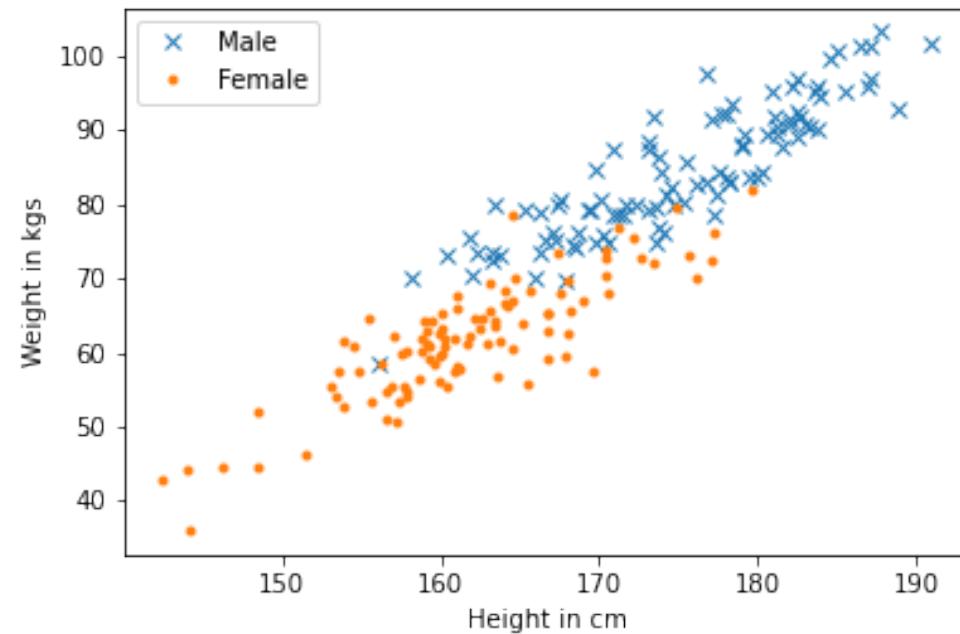
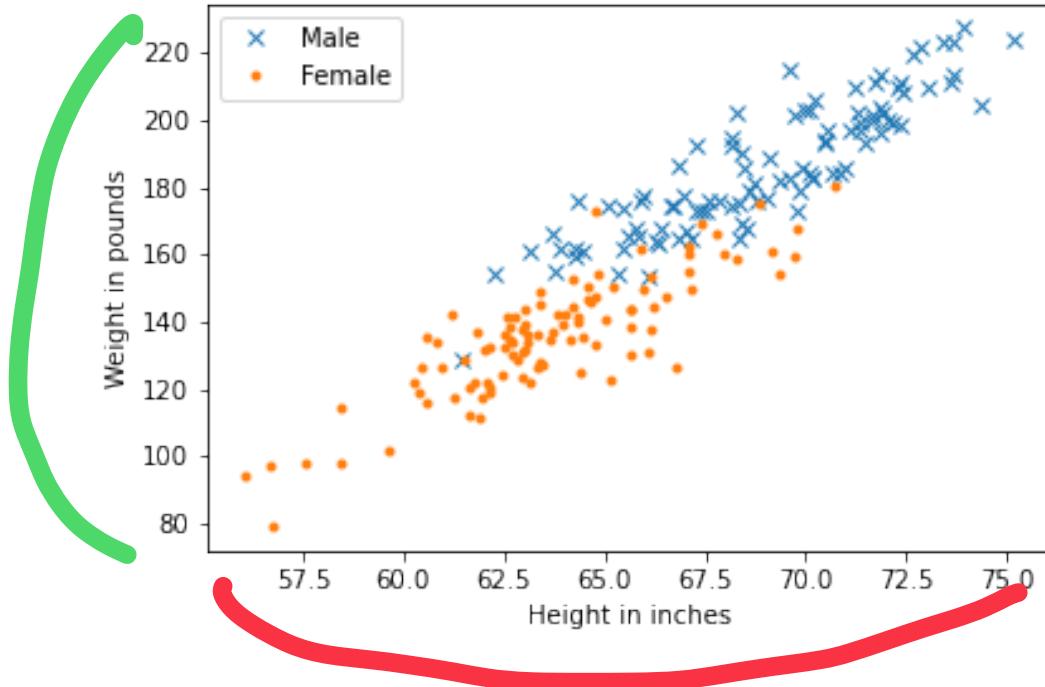
# Example: Height, weight, gender

- Dataset:
  - 10,000 observations
  - 5,000 of each gender
  - Height in inches
  - Weight in pounds
- <https://www.kaggle.com/mustafaali96/weight-height>
- Processing:
  - Shuffled
  - Split:
    - 5000 for training
    - 2500 for development testing
    - 2500 for final testing

Data format	Accuracy
Inch-pounds	0.905
Cm-kilos	0.897

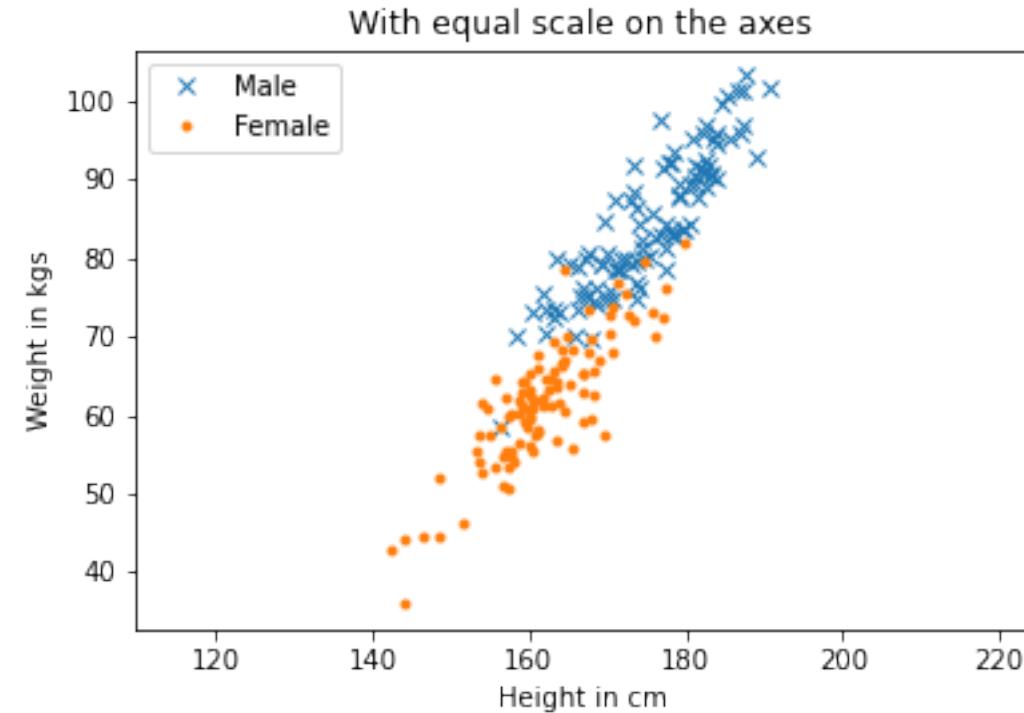
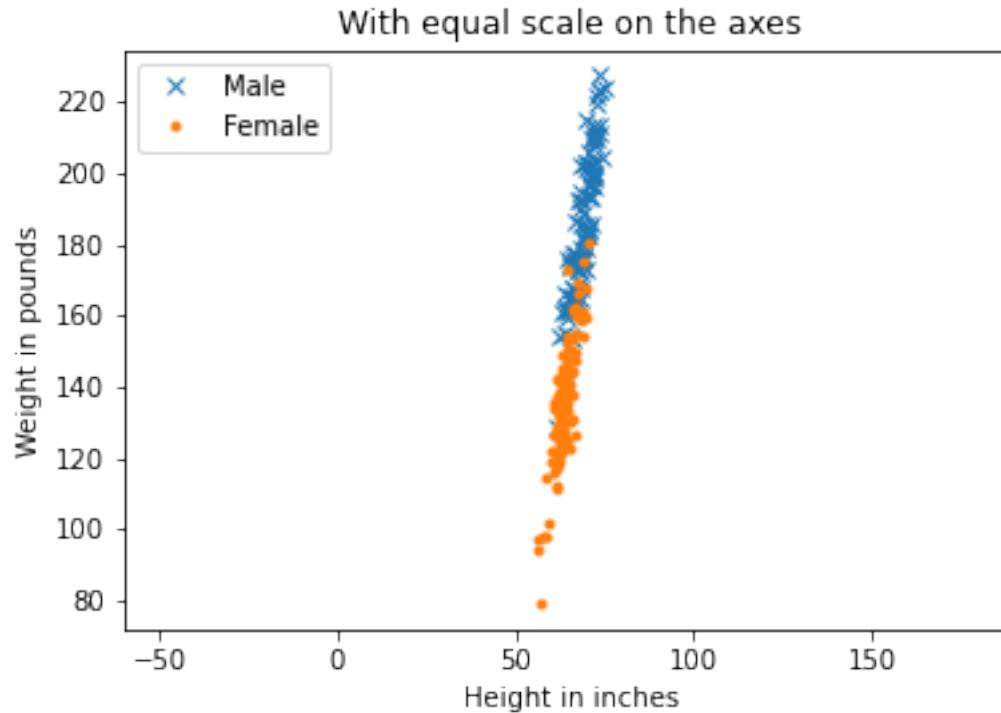
- We converted to cm and kgs using the same splits
- Trained and tested a 3NN-model
- Different results
- How come?

# How come?



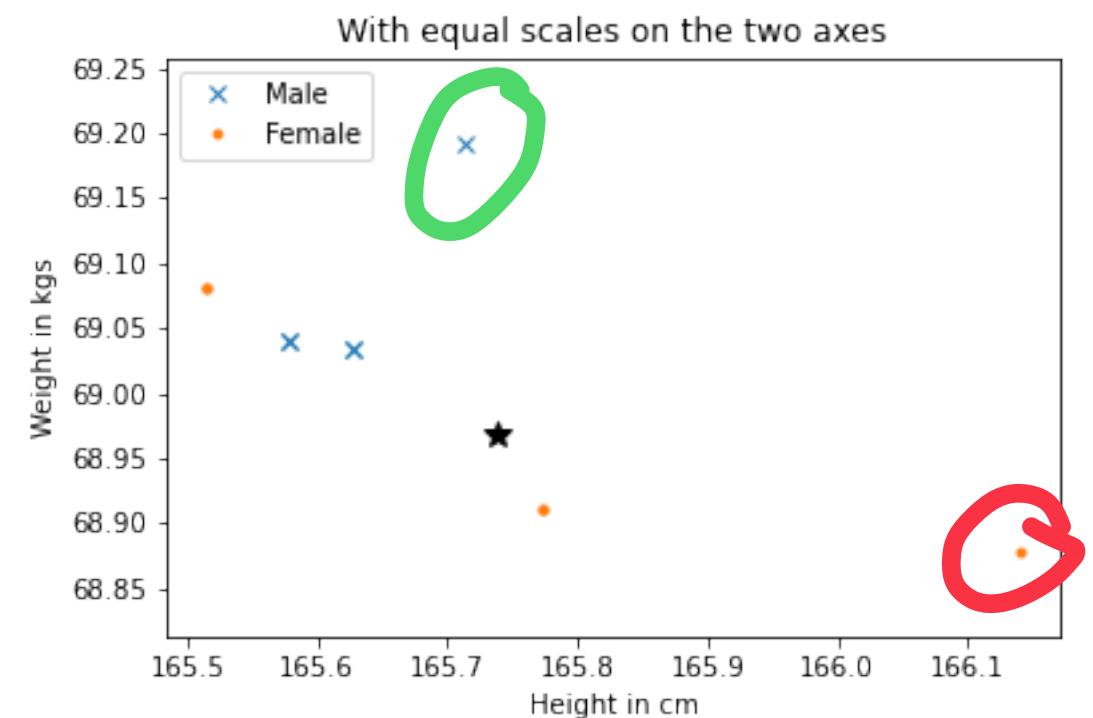
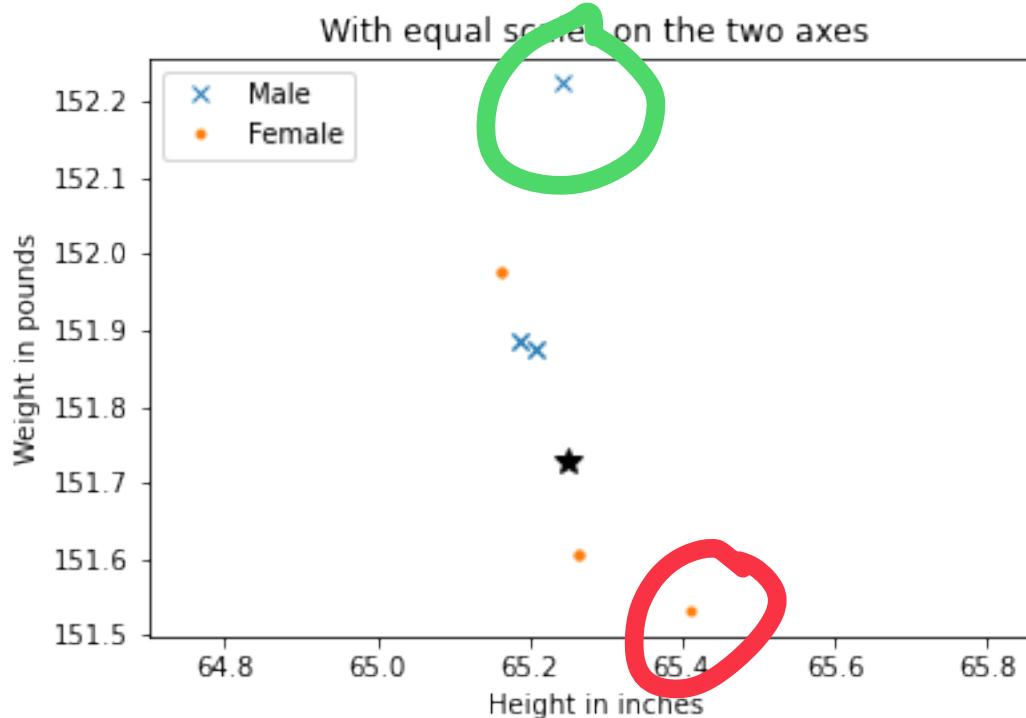
- The same random subsets of the training data
- Look the same

# Using equal scale at the two axes



- These are the spaces in which we calculate the distance to the nearest neighbors

# Using equal scale at the two axes



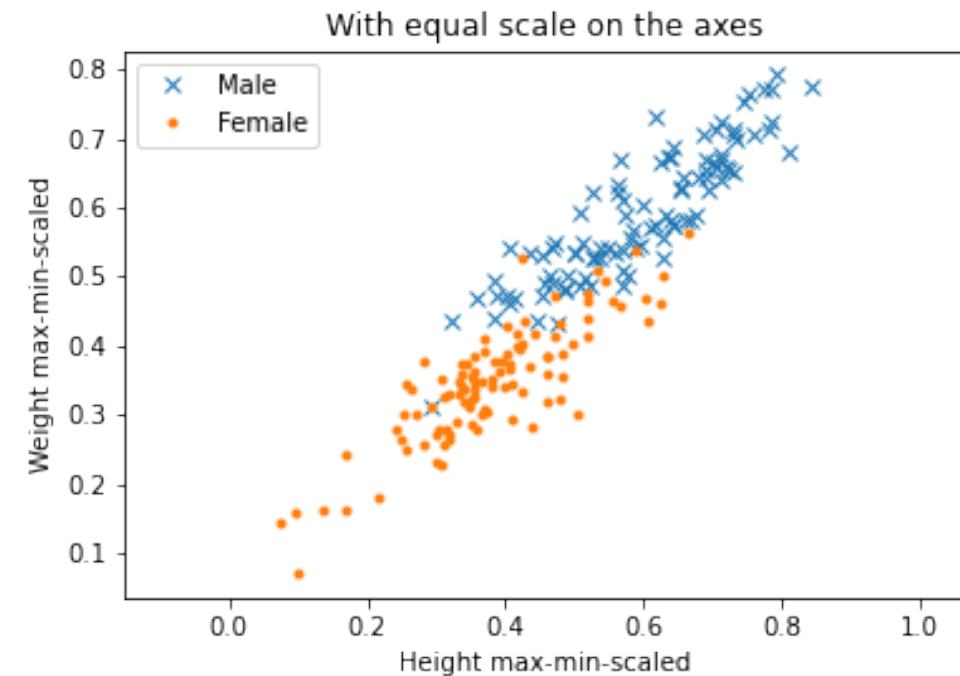
- Shows the same target and 6 training points in the two
- The 6 training points include the 5 nearest neighbors for each of them

# Scaling

- Observation:
  - The result of  $k$ NN depends on the scale we use on the feature dimensions.
- How do we determine the ``correct'' (i.e., best) scales?
  - We could experiment with various alternatives
- Another option is scaling
  - The input data should:
    - have similar range for the various input variables
    - typically, be in the range (0,1) or (-1, 1) or thereabouts
  - Several ways of scaling, most common
    - Max-min scaler
    - Normalization (standard scaler)

# Max-min-scaling

- Training set:  $N$  data points, of the form:
  - $\mathbf{x}_j = (x_{j,1}, x_{j,2}, \dots, x_{j,m})$
  - $m$  is the number of features
  - $x_{j,i}$  is the value of feature  $i$  for observation  $j$
- For each feature  $i = 1, \dots, m$ :
  - Let
    - $min_i = \min(\{x_{j,i} \mid j = 1, 2, \dots, N\})$
    - $max_i = \max(\{x_{j,i} \mid j = 1, 2, \dots, N\})$
  - Define  $scale_i(x_{j,i}) = \frac{x_{j,i} - min_i}{max_i - min_i}$
- This will map all features in the training set to numbers in the interval  $[0, 1]$



# Test set

- It is **not** a good idea to normalize the dataset before splitting it into training and testing (as Marsland will have it, p. 64):
  - Do not touch your test data
  - The goal is to apply the system also to new unseen data beyond the test data
- Instead:
  1. Use the training data to determine the scaler/normalizer
  2. Scale the training data before training
  3. Whenever you apply the system: scale the data before they are given to predict, **using the same scaler as on the training data**

# Normalizing (an alternative scaler)

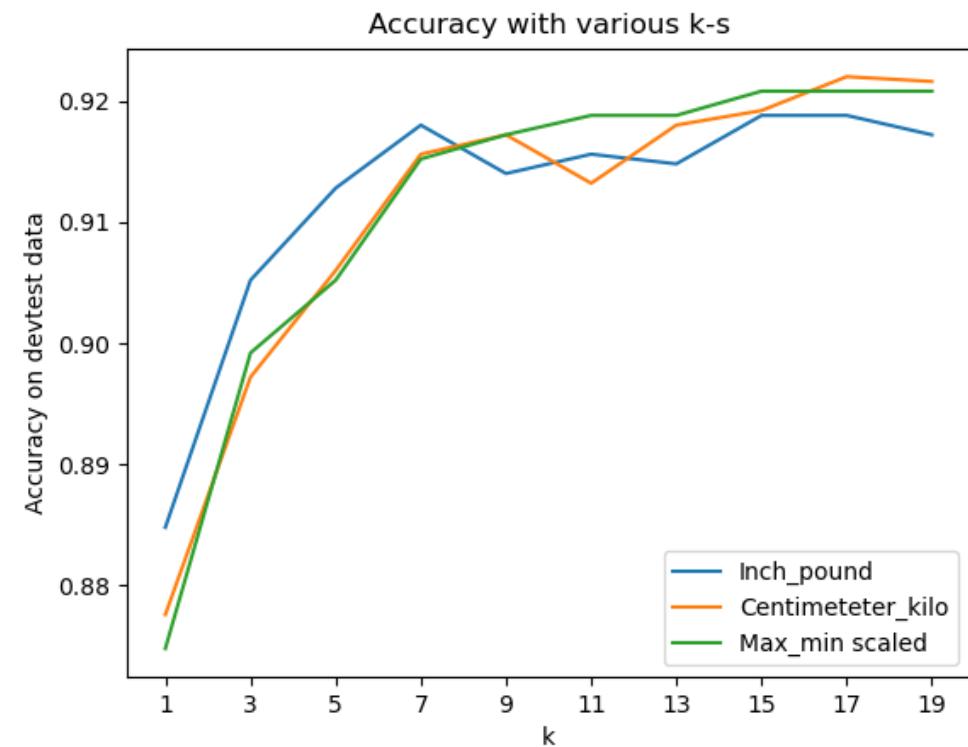
- Sometimes called standard scaler
- Find the
  - mean  $\mu_i = \text{mean}(\{x_{j,i} \mid j = 1, 2, \dots, N\})$ , and
  - standard deviation  $\sigma_i = \text{std}(\{x_{j,i} \mid j = 1, 2, \dots, N\})$
  - of your training set.
- Define  $scale_i(x_{j,i}) = \frac{x_{j,i} - \mu_i}{\sigma_i}$  for each feature  $i = 1, \dots, m$
- Apply it to your training data and to the data to which you apply the system.

# Scaling

- For some algorithms, scaling influences how items are classified, cf. *k*NN
- For other algorithms, the scaling
  - does not alter the result as the algorithm determine how much weight to put on various features
  - But scaling can be import for efficiency

# Results on the weight-height data

- In this case the scaler does not give better results.
- No guarantee for giving better results



# Take home

- The scale of the axes may influence the result of  $k\text{NN}$
- Scaling can be important for efficiency for some classifiers
- Max-min-scaling
- Standard scaler
- Construct scaler from training set alone

