


IN4050 - Introduction to Artificial Intelligence and Machine Learning

Multi-class Classification


Ali Ramezani-Kebrya





Multi-Class Classification One vs. Rest


IN4050 Introduction to Artificial Intelligence and Machine Learning



Multi-class classification

Classification

- Assign a label (class) from a finite set of labels to an observation



- So far, many algorithms and examples have been binary: *yes-no, 1-0*
- But many classification tasks are multi-class:
 - To each observation x choose one label from a finite set \mathbf{C}
 - What is different?

Multi-class classification

- A finite set, \mathbf{C} , of n different labels ($n > 2$)
- To each observation x choose one label from the set \mathbf{C}

We will consider two approaches:

- *One vs. rest classifier*
 - (also called one vs. all)
- *Multinomial logistic regression*, or *softmax regression*

1-of-N or "one hot encoding"

- The labels might be categorical:
 - 'apple', 'tomato', 'dog', 'horse'
 - The algorithms demand numerical attributes.
 - First attempt
 - 'apple' = 1
 - 'tomato' = 2
 - 'dog' = 3
 - etc.
 - Why isn't this a good idea?
- Better:
 - 'apple' = (1, 0, 0, 0, 0, 0)
 - 'tomato' = (0, 1, 0, 0, 0, 0)
 - 'dog' = (0, 0, 1, 0, 0, 0)
 - etc.
 - Both the target and the predicted value are vectors.

treats classification like regression

From multi-label to multi-class

Multi-label classifier


- Make n different classifiers, one for each class
- For classifier j :
 - consider class j the positive class
 - all other items in the negative class
 - train a classifier f_j
- Application
 - Assign a label c_j to an item if and only if it is classified as positive by f_j .

Multi-class classifier

- "To each observation x choose one label from the set C "
- How can a multi-label classifier be turned into a multi-class classifier?

One vs. rest (also called one vs. all)


- Start like the multi-label classifier: make one classifier for each class.
- It is easy to decide for items which fall into exactly one class
- But what if they fall into
 - More than one class?
 - No classes?



https://github.com/amueller/introduction_to_ml_with_python

One vs. rest

- If each classifier predicts a score, compare the scores for the classes
- Choose the class with the highest score.
- E.g., log. reg.:
 - Probability of being **red**: 0.8
 - Probability of being **blue**: 0.7
 - Choose **red**




https://github.com/amueller/introduction_to_ml_with_python



Multi-Class Classification Multinomial Logistic Regression

IN4050 Introduction to Artificial Intelligence and Machine Learning




Towards Multinomial Logistic Regression

- On the way, we will adopt a broader perspective
- Comparing perceptron, linear regression, logistic regression
- Compare to Marsland
- Shortly describe a multi-class perceptron

A general view


- Common
 - $z_j = \sum_{i=0}^n w_{i,j} x_i$
 - $w_{i,j}$ is the weight into node j from node i
 - Some index in opposite order

Binary classifiers		
Classifier	g	Loss
Perceptron	Step	0-1 loss
Lin. Regr.	Identity	MSE
Log.Regr.	Logistic	Cross-entropy




Connections going into a node

$$\begin{bmatrix} x_0 & x_1 & x_2 & \cdots & x_m \end{bmatrix} \begin{bmatrix} w_{0,1} & w_{0,2} & \cdots & w_{0,n} \\ w_{1,1} & w_{1,2} & \cdots & w_{1,n} \\ w_{2,1} & w_{2,2} & \cdots & w_{2,n} \\ \vdots & \vdots & \ddots & \vdots \\ w_{m,1} & w_{m,2} & \cdots & w_{m,n} \end{bmatrix} = \begin{bmatrix} z_1 & z_2 & \cdots & z_n \end{bmatrix}$$




Connections going out of a node

$$\begin{bmatrix} x_0 & \textcircled{x}_1 & x_2 & \cdots & x_m \end{bmatrix} \begin{bmatrix} w_{0,1} & w_{0,2} & \cdots & w_{0,n} \\ w_{1,1} & w_{1,2} & \cdots & w_{1,n} \\ w_{2,1} & w_{2,2} & \cdots & w_{2,n} \\ \vdots & \vdots & \ddots & \vdots \\ w_{m,1} & w_{m,2} & \cdots & w_{m,n} \end{bmatrix} = \begin{bmatrix} z_1 & z_2 & \cdots & z_n \end{bmatrix}$$




Multi-label perceptron

- Marsland's description of perceptron:
 - Possible targets: $(0,1,0,0,1)$
 - $y_j = g(z_j) = \begin{cases} 1 & \text{if } z_j > 0 \\ 0 & \text{if } z_j \leq 0 \end{cases}$
 - Loss is 0-1 loss for each j
- Describes a multi-label classifier
 - Each y_j depends only on the $w_{i,j}$'s
 - This could have been described as n independent binary perceptrons




Multi-class perceptron

- The target contains one 1, the rest are 0s
- $g(z_1, z_2, \dots, z_n) =$
- $\text{argmax} (z_1, z_2, \dots, z_n)$
 - The index with the max value
- The update rule (0-1 loss)
 - $w_{i,j} = w_{i,j} - \eta(y_j - t_j)x_i$
 - will correct for $j = 2$ and $j = 5$
 - leaves the other weights unaltered




Multi-output linear regression

- $y_j = g(z_j) = z_j$
- MSE-loss: $\sum_{k=1}^N (\sum_{j=1}^n (y_{k,j} - t_{k,j})^2)$
 - n output nodes
 - N input items
- y_j independent of $w_{i,k}$ $k \neq j$,
 - hence corresponds to n independent models
 - (Gets more interesting for multi-layer networks)




Multinomial Logistic Regression

- $z_j = \sum_{i=0}^n w_{i,j} x_i$
- Apply the softmax-function, S , where
 - $y_j = (S(z_1, \dots, z_n))_j = \frac{e^{z_j}}{\sum_{k=1}^n e^{z_k}}$
- Observe:
 - y_j depends on all the z_k
 - If $z_h > z_k$ then $y_h > y_k$
 - $0 < y_j < 1$
 - $\sum_{j=1}^n y_j = 1$
 - A probability distribution
 - $\hat{P}(C_j | \vec{x}) = \frac{e^{\vec{w}_j \cdot \vec{x}}}{\sum_{k=1}^n e^{\vec{w}_k \cdot \vec{x}}}$




Training Multinomial Logistic Regression 1

- The target has the form $(0,0,\dots,0,1,0,\dots,0)$, say
 - $t_s = 1$ and $t_j = 0$ for $j \neq s$
- We compare
 - $y = (y_1, y_2, \dots, y_n)$
- to the target labels
 - $t = (t_1, t_2, \dots, t_n)$
- using cross-entropy loss
 - $L_{CE}(y, t) = - \sum_{j=1}^n t_j \log y_j = - \log y_s$




Training Multinomial Logistic Regression 2

- $y_j = \frac{e^{z_j}}{\sum_{k=1}^n e^{z_k}}$
- $L_{CE}(\mathbf{y}, \mathbf{t}) = -\sum_{j=1}^n t_j \log y_j = -\log y_s$
- Goal: to find $\frac{\partial}{\partial w_{i,j}} L_{CE}(\mathbf{x}, \mathbf{t}, \mathbf{w})$ for all $w_{i,j}$
- Use the chain-rule for derivatives
- A little more complicated than for LogReg
- The result is simple, though
- $\frac{\partial}{\partial w_{i,j}} L_{CE}(\mathbf{x}, \mathbf{t}, \mathbf{w}) = (y_j - t_j)x_i$



Training Multinomial Logistic Regression 3

- $w_{i,j} = w_{i,j} + \eta(t_j - y_j)x_i$
- if $t_s = 1$:
 - $w_{i,s} = w_{i,s} + \eta(1 - y_s)x_i$
 - $w_{i,j} = w_{i,j} - \eta(y_j)x_i$
 - for $j \neq s$
- Here
 - $z_j = \sum_{i=0}^m w_{i,j}x_i$
 - $y_j = \frac{e^{z_j}}{\sum_{k=1}^n e^{z_k}}$
- Observe: All weights are updated for each observation



Applying Multinomial Logistic Regression

- We can use this as a probabilistic classifier

- $\hat{P}(C_j | \vec{x}) = \frac{e^{\vec{w}_j \cdot \vec{x}}}{\sum_{k=1}^n e^{\vec{w}_k \cdot \vec{x}}}$

- To make hard decisions use

$$\text{argmax}_{j=1, \dots, n} \frac{e^{\vec{w}_j \cdot \vec{x}}}{\sum_{k=1}^n e^{\vec{w}_k \cdot \vec{x}}}$$

