

IN3050/IN4050, Lecture 10

Unsupervised Learning

1: Introduction

2: Learning how to generate similar data

3: Learning how data are grouped

4: Learning how to compress data


5: Learning how to represent data in low dimension

IN3050/IN4050, Lecture 10

Unsupervised Learning

1: Introduction

Ole Christian Lingjærde



Supervised Learning


Unsupervised Learning

Supervised vs. unsupervised learning

Supervised learning: we have a training set consisting of a feature vector $\mathbf{x} = (x_1, \dots, x_m)$ and a response y for each sample. Goal is to learn the mapping $f: \mathbf{x} \mapsto y$.


Unsupervised learning: we have a training set consisting of only a feature vector $\mathbf{x} = (x_1, \dots, x_m)$ for each sample. Goal is to learn something interesting about the distribution of the \mathbf{x} 's.

Example (1D)




Each sample has a single feature x (a real number)

Example (2D)




Each sample has two features $\boldsymbol{x} = (x_1, x_2)$

Example (3D)



Each sample has three features $\mathbf{x} = (x_1, x_2, x_3)$


Example (p dimensions)



Each sample has 66 features $\mathbf{x} = (x_1, x_2, x_3, \dots, x_{66})$

What can we learn in unsupervised learning?

- **How to generate similar data**
→ Density estimation, generative models
- **How many (and which) groups are present**
→ hierarchical clustering, k-means clustering
- **How to compress the data**
→ Autoencoders
- **How to visualize the data in low dimension**
→ PCA, learning vector quantization, self organizing maps



Evaluating the performance

Supervised learning:

- Goodness of fit to training data
- Goodness of fit to new data
- Prediction performance



Unsupervised learning:

- No "truth" to compare with
- Performance more open to interpretation
- Still: often possible to define a useful loss function and to determine which of two solutions is best

IN3050/IN4050, Lecture 10

Unsupervised Learning

3: Learning how data are grouped
Ole Christian Lingjærde



Clustering

Learning how the data are grouped is called clustering.

Some clustering methods:

- Hierarchical clustering
- k-means clustering
- Competitive learning

Hierarchical clustering

- 1) Initially, each sample is a separate cluster
- 2) Find the two clusters that are closest together and merge them into one cluster
- 3) Repeat point 2 until there is only one cluster left




Hierarchical clustering: important aspects

- Must keep track of the whole history of cluster mergers in order to plot the tree (= dendrogram)
- In each iteration we have to compute the distance between every pair of the remaining clusters
- We have to define what we mean by the distance between two clusters


Hierarchical clustering: how to define the distance between two clusters C_1 and C_2

- Referred to as linkage method
- Has great influence on result
- Most common linkage methods:
 - Single linkage
 - Average linkage
 - Complete linkage
 - Ward linkage




Single linkage

$d(C_1, C_2) = \text{smallest distance between a sample in } C_1 \text{ and a sample in } C_2$




Average linkage

$d(C_1, C_2)$ = average distance between samples
in C_1 and samples in C_2




Complete linkage

$d(C_1, C_2) = \text{maximal distance between a sample in } C_1 \text{ and a sample in } C_2$



Ward linkage


$d(C_1, C_2)$ = increase in within-cluster variance if clusters are merged



Hierarchical clustering: how to define the distance between two samples \mathbf{x}_1 and \mathbf{x}_2

Manhattan distance:

$$d(\mathbf{x}_1, \mathbf{x}_2) = \sum_{i=1}^p |x_{1i} - x_{2i}|$$



Euclidean distance:

$$d(\mathbf{x}_1, \mathbf{x}_2) = \sqrt{\sum_{i=1}^p (x_{1i} - x_{2i})^2}$$

Cosine distance:


$$d(\mathbf{x}_1, \mathbf{x}_2) = 1 - \cos \theta = 1 - \frac{\mathbf{x}_1 \cdot \mathbf{x}_2}{\|\mathbf{x}_1\| \cdot \|\mathbf{x}_2\|}$$

k-means clustering

As before we wish to cluster samples $x_1, \dots, x_n \in \mathbf{R}^p$


- 1) Randomly select k cluster centers $c_1, \dots, c_k \in \mathbf{R}^p$
- 2) Assign each sample to the nearest cluster center
- 3) Update cluster centers to mean of samples in cluster
- 4) Repeat steps 2-3 until convergence

Example: k-means with k=3



Voronoi diagram

Voronoi cell: part of the plane that is closer to one particular cluster center than to any other cluster center




k-means clustering: algorithmic details

- Initial cluster center selection:
 - Random (uniform, normal) coordinates within data range
 - Random selection of the data points
 - Manual selection
- Distance between a sample x_i and a cluster center c_i :
$$d(x_i, c_i) = \sum (x_{ij} - c_{ij})^2 \quad (\text{or some other metric})$$
- Cluster centers are updated by calculating the average of all the sample vectors assigned to it:

$$c_i^{new} = \frac{1}{m} \sum x_i \quad (\text{or median})$$


The k-means neural network

- Normalize all samples ($\|x\|=1$ for all input vectors)
- Computing output values for input x :
 - Calculate $h_k = x \cdot w_k$ for all output nodes
 - Set activation value to 1 for output node with maximal h_k and set activation to 0 for all other output nodes
- Learning rule: $\Delta w_{ij} = \eta(x_j - w_{ij})$



k-means \neq k-means neural network

A: The normalization in k-means neural network means we lose information (only angular distance is preserved):



B: The learning rule in k-means neural network is on-line (i.e. weights are updated after each sample)